

9 Image Segmentation

9.1 Image Segmentation

9.2 Thresholding

9.3 Classification

9.4 K-means Clustering

9.5 Bayesian Classification

9.6 Region-Based Segmentation

9.7 Video Segmentation

9.1 Image Segmentation

Decomposition of scene into its **components**

- Key step in image analysis and object-based coding
- Correspondence to **physical objects** (ideal)



Spatial segmentation

- Into different regions

Temporal segmentation



- Shot detection

Requirements for complete segmentation

- **Connectivity**: each region (segment) consists of connected image points
- **Completeness**: union of all regions yields complete image
- **Homogeneity**: each region is homogeneous under given criterion
- **Closeness**: combining two segments gives inhomogeneous region

Result of segmentation process is also called **partition**

Partition Examples



Cluster-Based Segmentation

Assumptions


- Each pixel is assigned a set of **features** (color, gradient, texture, etc.)
- Feature space reveals significant cumulative **clusters**

① **Supervised** segmentation (**classification**)

- Class prototypes are known (e.g. pdfs)
a priori knowledge available

② **Unsupervised** segmentation (**cluster analysis**)

- Neither class prototypes nor number of classes are available



Example: 2 features, 3 classes

Special case: **thresholding**


- Only **one feature** (e.g. luminance)
- Only two classes: object (foreground) vs. background

9.2 Thresholding

Key issue: threshold θ selection

Supervised thresholding

- Example: MAP (Maximum a Posteriori) estimator
 - Minimizes segmentation error



Unsupervised Thresholding

Idea: find Θ that **minimizes within-class variance**

- Operates directly on **gray level** histogram
- Assumes **bimodal** distribution (foreground vs background)

(*)
$$\sigma_{wcv}^2(\Theta) = \underbrace{\omega_F(\Theta)}_{\text{prob. of foreground}} \sigma_F^2(\Theta) + \underbrace{\omega_B(\Theta)}_{\text{prob. of background}} \sigma_B^2(\Theta)$$
 where $\omega_{F/B}(\Theta) = \frac{N_{F/B}(\Theta)}{N}$
 $\omega_F(\Theta) + \omega_B(\Theta) = 1$

Issues

- Requires exhaustive search
- Computing variances can be computationally expensive

Total variance does not change

- Equivalent solution: **maximize between-class variance**

$$\sigma_{bcv}^2(\Theta) = \overset{\text{total var.}}{\sigma^2} - \sigma_{wcv}^2(\Theta)$$

Unsupervised Thresholding

Between-class variance

$$\begin{aligned}
 \sigma_{\text{bcv}}^2(\Theta) &= \sigma^2 - \sigma_{\text{wcv}}^2(\Theta) = \quad (\times) \\
 &= \left[\left(\frac{1}{N} \sum_n f^2[\mathbf{n}] \right) - \mu^2 \right] - \frac{N_F}{N} \left[\left(\frac{1}{N} \sum_{n \in F} f^2[\mathbf{n}] \right) - \mu_F^2 \right] - \frac{N_B}{N} \left[\left(\frac{1}{N} \sum_{n \in B} f^2[\mathbf{n}] \right) - \mu_B^2 \right] = \\
 &= -\mu^2 + \frac{N_F}{N} \mu_F^2 + \frac{N_B}{N} \mu_B^2 + \frac{1}{N} \left(\sum_n f^2[\mathbf{n}] - \frac{N_F}{N} \sum_{n \in F} f^2[\mathbf{n}] - \frac{N_B}{N} \sum_{n \in B} f^2[\mathbf{n}] \right) \quad 0
 \end{aligned}$$

$$\omega_F(\Theta) + \omega_B(\Theta) = 1$$

$$\omega_F(\Theta)\mu_F(\Theta) + \omega_B(\Theta)\mu_B(\Theta) = \mu$$

$$\sigma_{\text{bcv}}^2(\Theta) = \omega_F(\Theta)\omega_B(\Theta)(\mu_F(\Theta) - \mu_B(\Theta))^2$$

*only mean required,
no variance !*

Unsupervised Thresholding

Θ

Search for threshold that maximizes between-class variance

- This automatic threshold selection is known as **Otsu's algorithm**

Efficient recursive computation of $N_{F/B}$ and $\mu_{F/B}$:

$$N_F(\Theta + 1) = N_F(\Theta) + n_\Theta$$

$$N_B(\Theta + 1) = N_B(\Theta) - n_\Theta$$

$$\mu_F(\Theta + 1) = \frac{\mu_F(\Theta)N_F(\Theta) + \Theta n_\Theta}{N_F(\Theta + 1)}$$

$$\mu_B(\Theta + 1) = \frac{\mu_B(\Theta)N_B(\Theta) - \Theta n_\Theta}{N_B(\Theta + 1)}$$

calculate σ_{bc}^2 using (*)
→ select max.

(*)


n_Θ - Θ^{th} bin of image histogram (number of pixels with luminance equal to Θ)

Otsu's Thresholding Examples

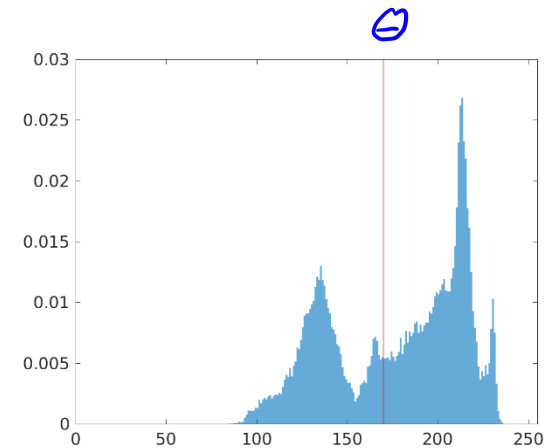

Original image



Thresholded image




Histograms




=> use morphological filters for smoothing binary mask !

Otsu's Thresholding Examples

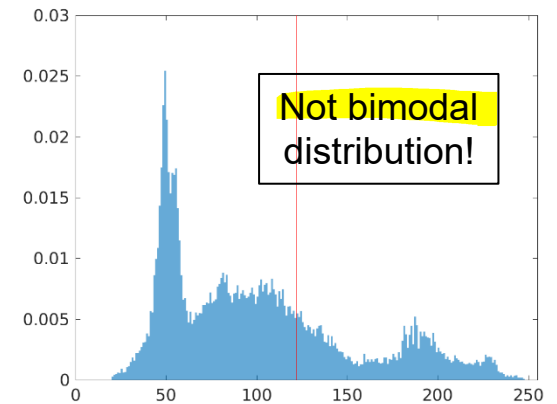

Original image



Thresholded image



Histograms




9.3 Classification

For each pixel (m, n) we have L features

- Feature vector $\mathbf{f}[m, n] = [f_0[m, n], f_1[m, n], \dots, f_{L-1}[m, n]]^T$

2D example, $L=2$



We have K clusters (a priori knowledge)

- Cluster centroids $\mathbf{c}^{(k)} = [c^{(0)}, c^{(1)}, \dots, c^{(K-1)}]^T$ (k^{th} cluster)

Each pixel is assigned the **best fitting** cluster S according to

$$S[m, n] = \underset{k}{\operatorname{argmin}} \sum_{l=1}^L |f_l[m, n] - c_l^{(k)}|^P = \underset{k}{\operatorname{argmin}} \|\mathbf{f}[m, n] - \mathbf{c}^{(k)}\|_P$$

- If $P=2$ (Euclidean norm) we have nearest neighbor classification

Chroma Keying

Color is more powerful feature than luminance

- 3D space vs. 1D space

Classification with only **two classes**

- Take picture in front of green screen (or blue, or orange...)

movie "Matrix"



Taken picture



Classification



Merging



<http://techteacherslog.net>


<http://www.ralph-dte.eu>

<https://i.ytimg.com/vi/4bkENVYNHHs/maxresdefault.jpg>

9.4 K-means **Clustering** (unsupervised case)

Simple **unsupervised** learning algorithm

- Extension of nearest neighbor classification with **unknown** cluster centroids (clustering problem)



Assumption: number of classes K known a priori

Start with arbitrarily chosen set of K cluster centroids $\mathbf{c}^{(k)}$

- Step 1
- 1) **Nearest neighbor classification** (assign each pixel to its nearest cluster)
 - 2) Re-compute cluster centroids

$$\mathbf{c}_{\text{new}}^{(k)} = \frac{1}{N_k} \sum_{\substack{(m,n) \in \\ \text{cluster } k}} \mathbf{f}[m,n] \quad k = 0, 1, \dots, K - 1$$

Mean value of all features
assigned to class k

N_k – number of feature vectors
assigned to cluster k

- 3) Go back to step 1) (or terminate if centroids don't change anymore)

d. LBG in VQ [IVC lecture]

K-means Clustering Example

$l=1$, only grey level

Original





2 clusters



3 clusters



4 clusters



K-means Clustering

Automatically finds clusters that minimize squared classification error

- Important for **image quantization** (cluster centroids)

Example (8-bit luminance depth, i.e. 256 gray levels)

Original






(2 bits)
4 clusters



(3 bits)
8 clusters



(4 bits)
16 clusters



9.5 Bayesian Classification

Based on **minimization** of Bayes **risk** $R(\hat{S})$

- Risk defined as expected cost value

$$R(\hat{S}) = E[C(\hat{S}, S)] = \iint_{S, f} C(\hat{S}, S) P(S, f) dS df = \iint_{S, f} C(\hat{S}, S) \underbrace{P(S|f)P(f)}_{\text{Bayes rule}} dS df$$

estimated segmentation \nearrow \nearrow *true segmentation*

Feature probability $P(f)$ (“observed signal”) is class-independent

- The same for all classes \rightarrow no influence on minimization

$$R(\hat{S}|f) = \int_S C(\hat{S}, S) P(S|f) dS$$

Bayesian classification

$$\hat{S} = \operatorname{argmin}_S R(\hat{S}|f) = \operatorname{argmin}_S \left[\int_S C(\hat{S}, S) P(S|f) dS \right] \quad (*)$$

Maximum a Posteriori (MAP)

Cost function: $C(\hat{S}, S) = 1 - \delta(\hat{S}, S)$ $\delta(\hat{S}, S) = \begin{cases} 1 & \text{if } \hat{S} = S \\ 0 & \text{else} \end{cases}$ $c(\hat{S}, S) = \begin{cases} 0 & \text{if } \hat{S} = S \\ 1 & \text{if } \hat{S} \neq S \end{cases}$

Bayes risk: $R_{\text{MAP}}(\hat{S}|\mathbf{f}) = \int_S (1 - \delta(\hat{S}, S)) P(S|\mathbf{f}) dS = 1 - \underbrace{P(\hat{S}|\mathbf{f})}_{\rightarrow \max} \rightarrow \min$

Probability a posteriori $P(\hat{S}|\mathbf{f})$

- Probability of assigning \hat{S} given feature \mathbf{f} (usually not known)

Bayes' rule:

$$P(S|\mathbf{f}) = \frac{P(\mathbf{f}|S)P(S)}{p(\mathbf{f})}$$

does not affect maximization



Thomas Bayes
(1701-1761)

MAP classification:

$$\hat{S}_{\text{MAP}} = \underset{S}{\operatorname{argmax}} P(S|\mathbf{f}) = \underset{S}{\operatorname{argmax}} [P(\mathbf{f}|S)P(S)]$$

Bayesian Classification


Difference between MAP and nearest neighbor classification

- Decision not necessarily for the nearest cluster

Instead, the following is considered

- ① • Clusters a priori probability ($P(S)$)
 - Which cluster is more likely?
- ② • Feature likelihood within a specific class $P(f|S)$
 - Given a cluster, which feature vector is more likely?

Example (1D)



Nearest neighbor assigns class A (it is closer)

MAP assigns class B (it is more likely)

Bayesian Classification

If **classes are equally probable** then MAP is reduced to **Maximum Likelihood (ML)** classifier

$$\hat{S}_{\text{ML}} = \underset{S}{\operatorname{argmax}} [P(\mathbf{f}|S)P(S)] = \underset{S}{\operatorname{argmax}} P(\mathbf{f}|S)$$

constant

Other cost functions are also possible



Disadvantage of classification and clustering methods discussed so far

- Operate on features only (**no spatial relation** between pixels considered)

9.6 Region-Based Segmentation

Idea: Incorporate knowledge about **topological structure** of partition (especially **neighbor** relations)

Region: Group of connected pixels with similar properties
(=> *still incorporate features*)


Principles:

- ① Similarity
 - Feature differences / variance
- ② Spatial proximity *new here!*
 - Euclidean distance
 - Compactness of a region

Region Growing

Starts with a set of **seed** points


- Expansion (growth) by pixels with similar features
- Adjacent similar regions are merged together



→ bottom-up segmentation

Properties

- Fast and conceptually simple
- Sensitive to noise
- Gradient problem



Similarity Measures for Two Regions

Given mean μ and variance σ^2 of two neighboring regions R_i and R_j

① Absolute deviation of mean value

$$d(R_i, R_j) = |\mu(R_i) - \mu(R_j)|$$

- Simple to calculate, does not account for region variances

② Variance coherence

$$d(R_i, R_j) = \sigma^2(R_i \cup R_j) - \frac{\sigma^2(R_i) + \sigma^2(R_j)}{2}$$

(μ is included implicitly)

- Extensible to higher order statistics

③ Likelihood ratio

$$d(R_i, R_j) = \frac{\sigma^2(R_i \cup R_j)^{N_i + N_j}}{\sigma^2(R_i)^{N_i} \cdot \sigma^2(R_j)^{N_j}}$$

(includes reliability of estimates)



- Considers region **size N**

Region Splitting

Split image into disjoint regions

- Check each region for homogeneity, if not homogeneous keep splitting

Example: **quad-tree** decomposition



→ top-down segmentation

Problems:

- How to optimally split a region into homogeneous sub-regions?
- Requires knowledge about number of sub-regions and location of region boundaries (e.g. by **edge detection**)

Split & Merge Segmentation


Combination of agglomerative and divisive region operations

Splitting (e.g. quad-tree)

- Keep splitting until all blocks fulfill homogeneity criterion

Merging

- Merge all neighboring block which are sufficiently similar



Disadvantage: region borders often exhibit “staircase” character

9.5 Temporal Segmentation of Video

Detection of **scene cuts** and smooth scene transitions (fading)

- Usually as preprocessing for video structuring applications (news server, media abstraction for messaging)

Scene cut assumption

- Content changes significantly between two consecutive frames or over a number of frames



Analysis of suitable (global) **image feature** over time

- Color distribution, motion, texture descriptors, etc.
- Often audio track is analyzed in parallel to improve segmentation results (e.g. silence detection)


Shot Detection Using Feature Histograms

Shot detection

- Detect scene cuts and segment video into a number of temporal **consistent** video sequences (**shots**)

Histogram of suitable image feature (e.g. color)

- Analyze sum of absolute histogram differences between subsequent video frames
- Scene cut detected if: a) Sum exceeds hard threshold T_H (hard cut)
b) Remains above soft threshold T_S for a certain number of frames n_0 (soft cut or fading)



Sensitive to (sudden) illumination changes → use additional features


Shot Detection Using Motion Analysis

Change in **global motion** usually corresponds to new scene

- E.g. start and end of camera pan (= horizontal movement)

Detection of shot boundaries

- Significant change in **motion trajectory**



Can be improved by considering higher order motion

- Rotation, zoom, affine models, etc.