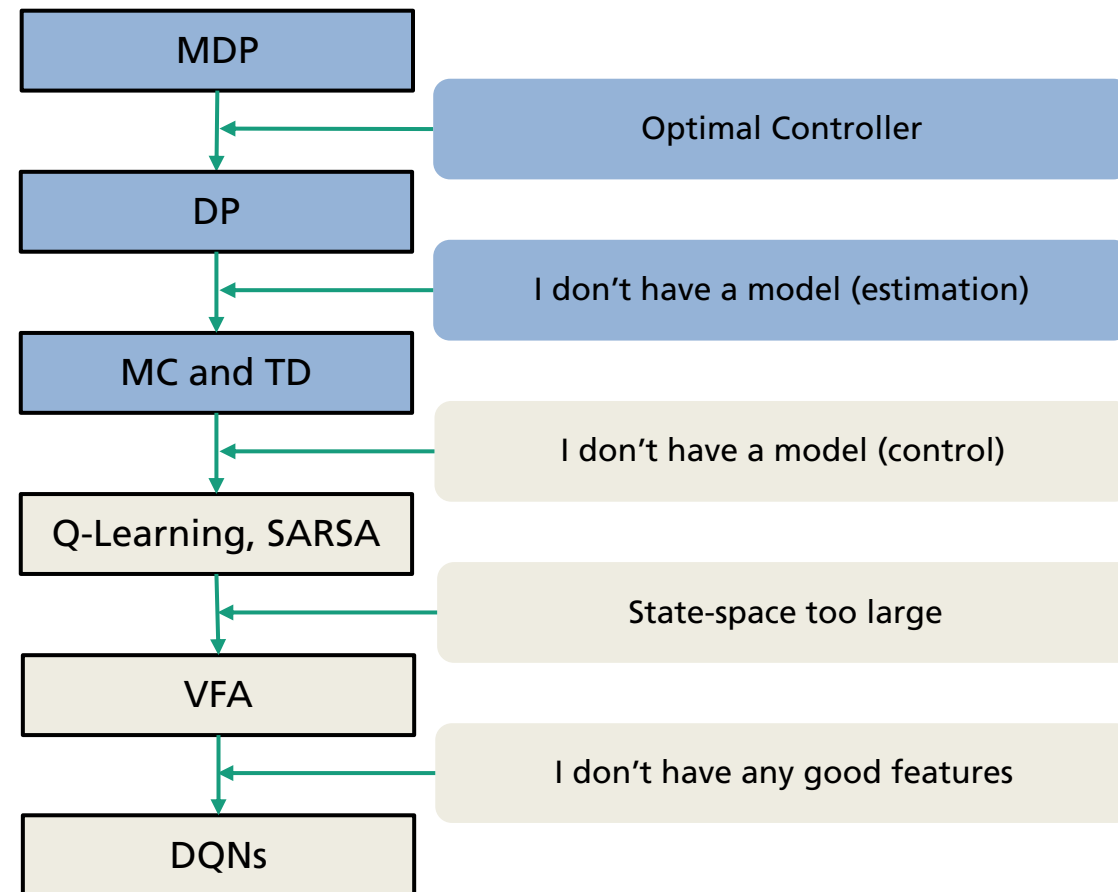


Model-free Prediction: TD-Learning

Christopher Mutschler



Overview



Monte Carlo and TD Methods

Assumptions:

- We know that the model of the world can be described by an MDP:

$$(\mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma)$$

- We know the (discrete) state and action spaces, i.e., \mathcal{S} and \mathcal{A} .
- We can interact with the world (with some policy π).
- We receive experience samples from the environment in the form

$$(S_t, A_t, R_t, S_{t+1}) = (s, a, r, s').$$

Monte Carlo and TD Methods

- Temporal-Difference Learning
 - Breaks up episodes and makes use of the intermediate returns
 - Learns directly from experience and interaction with the environment
 - Model-free: no knowledge of MDP
 - Learns from incomplete episodes (bootstrapping)
 - **We update a guess towards a guess**

Monte Carlo and TD Methods

- Temporal-Difference Learning: Idea of TD(0) Policy Evaluation

$$V^\pi(s) = \underbrace{r(s, \pi(s))} + \underbrace{\gamma \sum_{s' \in S} \mathcal{P}(s'|s, \pi(s)) V^\pi(s')}_{\text{We don't know the transition model}}$$

We don't know the transition model

$$(s, a, r, s')$$

But we have real transitions available

$$V^\pi(s) = r + \gamma V^\pi(s')$$

Let's assume that the reality is the transition we observed

$$V(s) \leftarrow V(s) + \alpha(r + \gamma V(s') - V(s))$$

→ and update our old estimate "a bit" in this direction

Monte Carlo and TD Methods

- TD(0) vs. MC Policy Evaluation
 - Goal: learn value function v_π online from experience when we follow policy π

- Simplest TD learning algorithm: TD(0)
- Update value **towards estimation \hat{G}** :

$$V(s) \leftarrow V(s) + \alpha(\hat{G} - V(s))$$

$$\hat{G} = r + \gamma V(s') \text{ (estimated return)}$$

- \hat{G} is called the TD target
- $\hat{G} - V(s)$ is called the TD error.

- Update $V(s)$ incrementally after each episode.
- For each state s with **actual return G** :

$$N(s) \leftarrow N(s) + 1 \quad \text{(just increment visit counter)}$$

$$V(s) \leftarrow V(s) + \frac{1}{N(s)} (G - V(s)) \quad \text{(update a bit } \rightarrow \text{ reduce error)}$$

- In non-stationary problems, it can be useful to track a running mean, i.e., forget old episodes:

$$V(s) \leftarrow V(s) + \alpha(G - V(s)).$$

$$\text{NewEstimate} \leftarrow \text{OldEstimate} + \text{StepSize} [\text{Target} - \text{OldEstimate}]$$

Monte Carlo and TD Methods

- TD(0) vs. MC Policy Evaluation

Tabular TD(0) for estimating v_π

Input: the policy π to be evaluated
 Algorithm parameter: step size $\alpha \in (0, 1]$
 Initialize $V(s)$, for all $s \in \mathcal{S}^+$, arbitrarily except that $V(\text{terminal}) = 0$
 Loop for each episode:
 Initialize S
 Loop for each step of episode:
 $A \leftarrow$ action given by π for S
 Take action A , observe R, S'
 $V(S) \leftarrow V(S) + \alpha[R + \gamma V(S') - V(S)]$
 $S \leftarrow S'$
 until S is terminal

First-visit MC prediction, for estimating $V \approx v_\pi$

Input: a policy π to be evaluated
 Initialize:
 $V(s) \in \mathbb{R}$, arbitrarily, for all $s \in \mathcal{S}$
 $Returns(s) \leftarrow$ an empty list, for all $s \in \mathcal{S}$
 Loop forever (for each episode):
 Generate an episode following π : $S_0, A_0, R_1, S_1, A_1, R_2, \dots, S_{T-1}, A_{T-1}, R_T$
 $G \leftarrow 0$
 Loop for each step of episode, $t = T-1, T-2, \dots, 0$:
 $G \leftarrow \gamma G + R_{t+1}$
 Unless S_t appears in S_0, S_1, \dots, S_{t-1} :
 Append G to $Returns(S_t)$
 $V(S_t) \leftarrow \text{average}(Returns(S_t))$


Sutton, R. S., & Barto, A. G. (2018). *Reinforcement learning: An introduction*. MIT press.

Monte Carlo and TD Methods

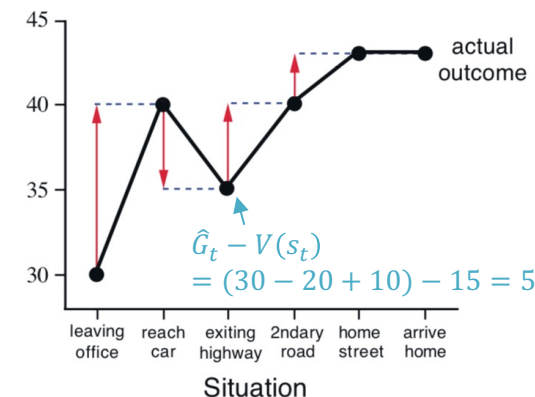
- Which one should I use? Does it make any difference?
- Example: Driving Home from work

State	Elapsed Time [min]	Predicted Time to Go [min]	Predicted Total Time [min]
leaving office, friday at 6	0	30	30
reach car, raining	5	35	40
exiting highway	20	15	35
2ndary road, behind truck	30	10	40
entering home street	40	3	43
arrive home	43	0	43

MC ($\alpha = 1$)





TD ($\alpha = 1$)




Sutton, R. S., & Barto, A. G. (2018). Reinforcement learning: An introduction. MIT press.

Monte Carlo and TD Methods

- Which one should I use? Does it make any difference?
- Example: Random Walk




RMS error,
averaged
over states



Sutton, R. S., & Barto, A. G. (2018). *Reinforcement learning: An introduction*. MIT press.

Monte Carlo and TD Methods


- Which one should I use? Does it make any difference?
 - TD can learn before (or even without) knowing the final outcome
 - after each step
 - incomplete sequences
 - continuing problems, very delayed or no return
 - MC only works for episodic problems (i.e., that terminate)
 - must wait until end of the episode



Sutton, R. S., & Barto, A. G. (2018). *Reinforcement learning: An introduction*. MIT press.

Monte Carlo and TD Methods

- Which one should I use? Does it make any difference?
 - Bias/Variance Trade-Off
 - MC has high variance, but zero bias
 - Good convergence (even with FA)
 - insensitive to initialization (no bootstrapping), simple to understand
 - TD has low variance, but some bias
 - TD(0) converges to $\pi_v(s)$ (be careful with FA: bias is a risk)
 - sensitive to initialization (because of the bootstrapping)
 - Usually more efficient in practice



Sutton, R. S., & Barto, A. G. (2018). *Reinforcement learning: An introduction*. MIT press.

Monte Carlo and TD Methods

- Which one should I use? Does it make any difference?
- Example: You are the predictor!
 - Two states A, B; no discounting; 8 episodes of experience
 - keep iterating on experience (MC and TD until both of them converge):

A, 0, B, 0

B, 1

B, 1


B, 1

B, 1

B, 1

B, 1

B, 0



- What is $V(S = A)$ and $V(S = B)$?


• MC: $V(A) = 0$ $V(B) = 0.75$

• TD: $V(A) = 0.75$ $V(B) = 0.75$

Sutton, R. S., & Barto, A. G. (2018). Reinforcement learning: An introduction. MIT press.

Monte Carlo and TD Methods

- Which one should I use? Does it make any difference?
 - TD exploits Markov property and is more efficient in Markov environments
 - MC is more efficient in non-Markov environments
 - TD usually converges faster than MC



Sutton, R. S., & Barto, A. G. (2018). *Reinforcement learning: An introduction*. MIT press.

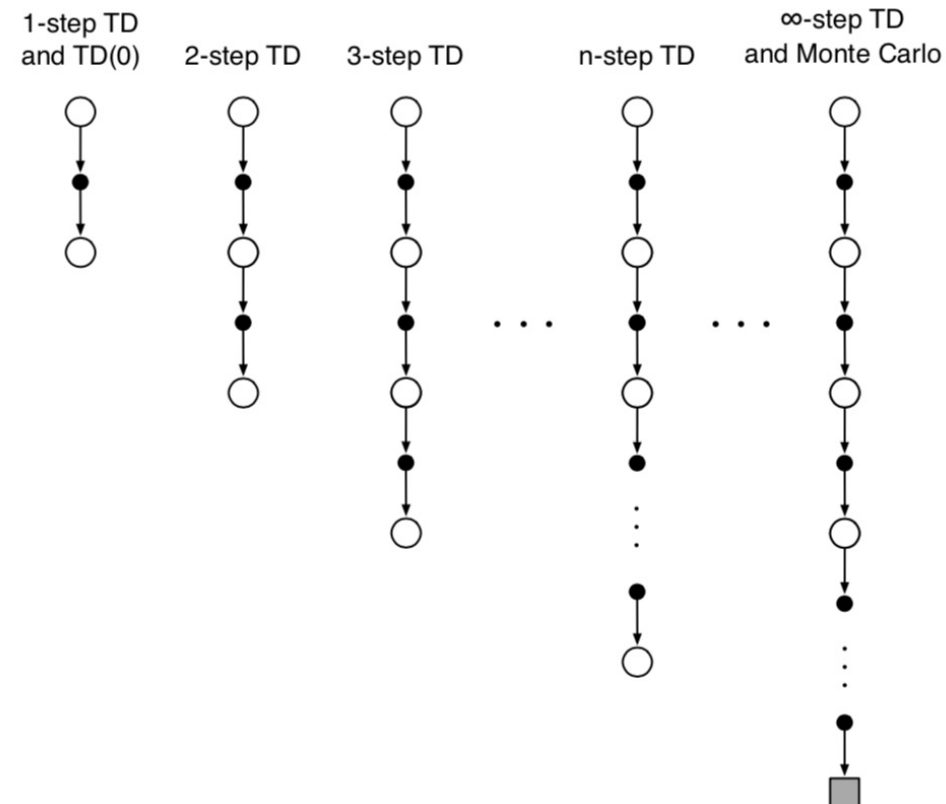
Monte Carlo and TD Methods

Hands-On:

https://cs.stanford.edu/people/karpathy/reinforcejs/gridworld_td.html

Monte Carlo and TD Methods

- Intermediate methods between MC and TD(0) exist
- They are based on n-step returns



Sutton, R. S., & Barto, A. G. (2018). *Reinforcement learning: An introduction*. MIT press.

Monte Carlo and TD Methods

- Intermediate methods between MC and TD(0) exist
- They are based on n-step returns

n-step TD for estimating $V \approx v_\pi$

```


Input: a policy  $\pi$ 
Algorithm parameters: step size  $\alpha \in (0, 1]$ , a positive integer  $n$ 
Initialize  $V(s)$  arbitrarily, for all  $s \in \mathcal{S}$ 
All store and access operations (for  $S_t$  and  $R_t$ ) can take their index mod  $n + 1$ 

Loop for each episode:
  Initialize and store  $S_0 \neq \text{terminal}$ 
   $T \leftarrow \infty$ 
  Loop for  $t = 0, 1, 2, \dots$ :
    If  $t < T$ , then:
      Take an action according to  $\pi(\cdot | S_t)$ 
      Observe and store the next reward as  $R_{t+1}$  and the next state as  $S_{t+1}$ 
      If  $S_{t+1}$  is terminal, then  $T \leftarrow t + 1$ 
       $\tau \leftarrow t - n + 1$  ( $\tau$  is the time whose state's estimate is being updated)
      If  $\tau \geq 0$ :
         $G \leftarrow \sum_{i=\tau+1}^{\min(\tau+n, T)} \gamma^{i-\tau-1} R_i$ 
        If  $\tau + n < T$ , then:  $G \leftarrow G + \gamma^n V(S_{\tau+n})$  ( $G_{\tau:\tau+n}$ )
         $V(S_\tau) \leftarrow V(S_\tau) + \alpha [G - V(S_\tau)]$ 
  Until  $\tau = T - 1$ 
  
```

Sutton, R. S., & Barto, A. G. (2018). Reinforcement learning: An introduction. MIT press.

Monte Carlo and TD Methods


- Intermediate methods between MC and TD(0) exist
- They are based on n-step returns
- Unfortunately, their prediction accuracy is sensitive to the algorithm hyperparameters
- Example: Random Walk



Sutton, R. S., & Barto, A. G. (2018). Reinforcement learning: An introduction. MIT press.

Monte Carlo and TD Methods

- A Unified View of Prediction Algorithms



Sutton, R. S., & Barto, A. G. (2018). *Reinforcement learning: An introduction*. MIT press.

TD – Remarks on Convergence Properties

- There is a lot of work that studied the convergence of TD:
- Convergence and optimality of (linear) TD methods under batch training (no online learning):
 - Richard S. Sutton: Learning to predict by the Methods of Temporal Differences. Machine Learning 3:9-44. 1988.
- Build on [Sutton1988] and proofs convergence of TD(0) and extends Watkin's Q-learning theorem (next video):
 - Peter Dayan: The Convergence of TD(λ) for General λ . Machine Learning 8:341-362. 1992.
- Further studies in the context of Q-Learning and SARSA (next video):
 - Tommi Jaakkola, Michael Jordan, Stander Singh: On the Convergence of Stochastic Iterative Dynamic Programming Algorithms. Technical report. 1994.
 - Francisco Melo: Convergence of Q-Learning: A Simple Proof. Technical report. (it has only 4 pages – so feel free to have a look 😊)
 - Satinder Singh, Tommi Jakkola, Michael Littman, Csaba Szepesvari: Convergence Results for Single-Step On-Policy Reinforcement-Learning Algorithms. Machine Learning 39:287-308. 2000.