

Policy-based Reinforcement Learning

Christopher Mutschler



Vapnik's rule

“Never solve a more general problem as an intermediate step.”

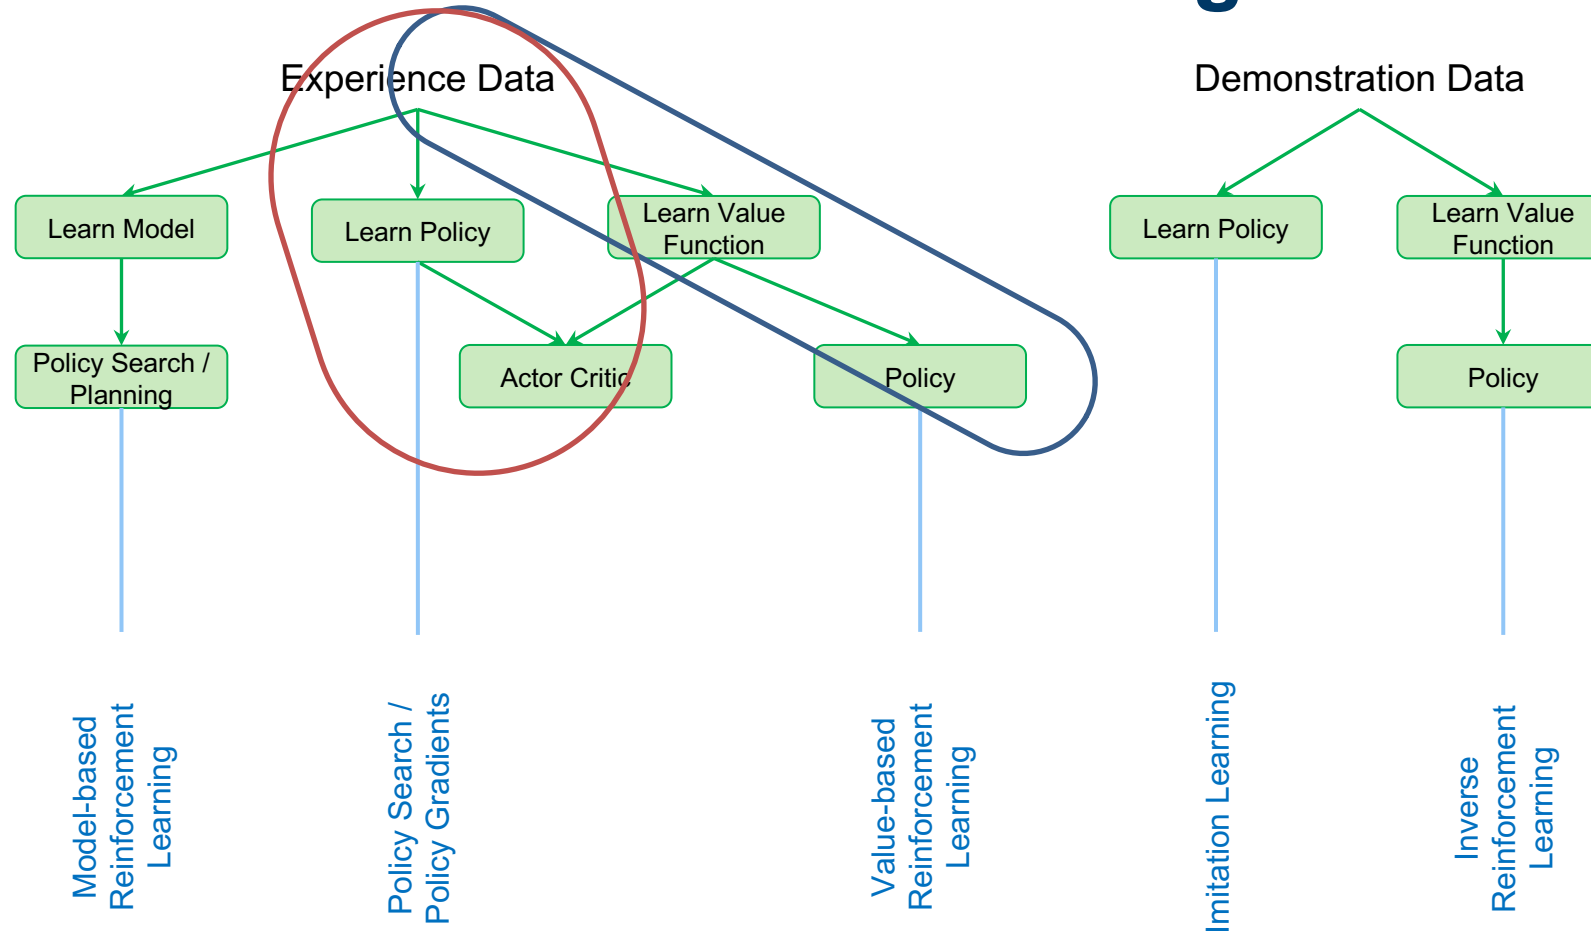
- Vladimir Vapnik, 1998

- Remember:

”New goal: find a policy that maximizes the expected return!”

- If we care about optimal behavior: why not learn a policy directly?

Policy-based Reinforcement Learning

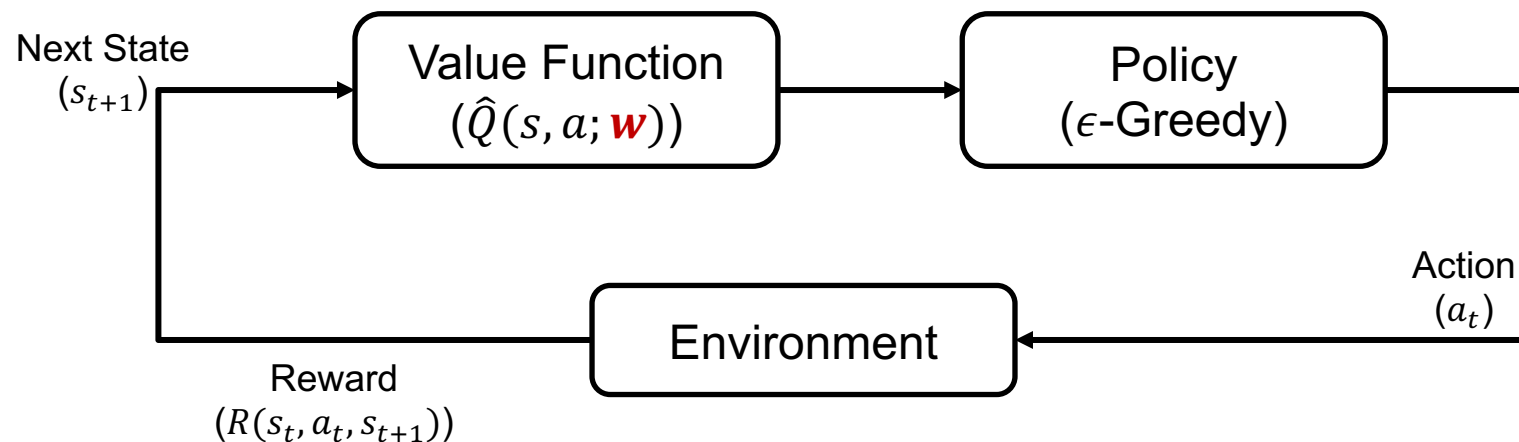


Policy-based Reinforcement Learning

- Previously we approximated parametric value functions:

$$v_w(s) \approx v_\pi(s)$$
$$q_w(s, a) \approx q_\pi(s, a)$$

- A policy can be generated from these values
 - e.g., greedy or ϵ -greedy



Goal: find w that approximates the true Q -function

Policy-based Reinforcement Learning

- Previously we approximated parametric value functions:


$$v_w(s) \approx v_\pi(s)$$

$$q_w(s, a) \approx q_\pi(s, a)$$

- A policy can be generated from these values
- In this lesson we will directly parameterize the policy:

$$\pi_\theta(a|s) = p(a|s; \theta)$$

- We still focus on model-free reinforcement learning



Goal: find θ that maximizes long term reward

General overview

- **Model-based RL:**


- + “Easy” to learn a model (supervised learning)
- + Learns *all there is to know* from the data
- Objective captures irrelevant information
- May focus computations/capacity on irrelevant details
- Computing policy (planning) is non-trivial and can be computationally expensive

- **Value-based RL:**

- + Closer to true objective
- + Fairly well-understood: somewhat similar to regression
- Still not the true objective: may still focus capacity on less-important details

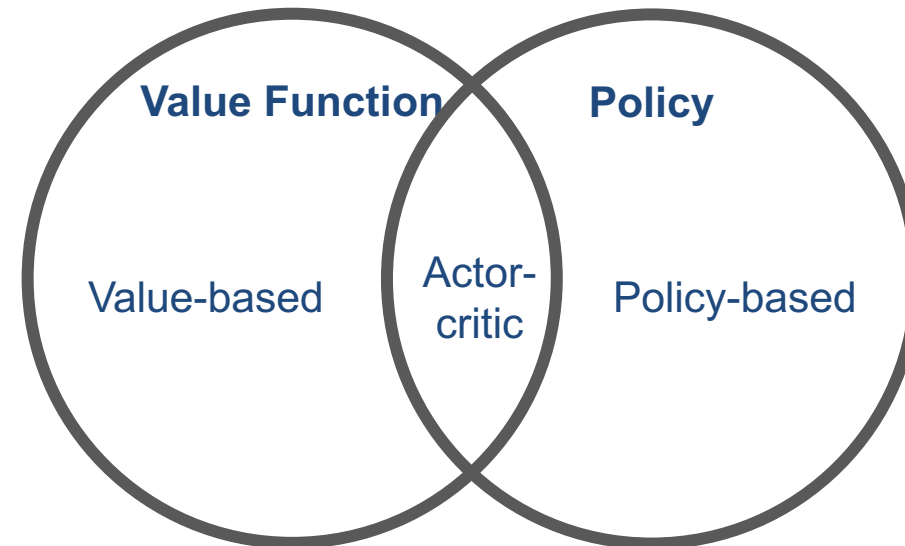
- **Policy-based RL:**

- + Right objective!
- Ignores other learnable knowledge (potentially not the most efficient use of data)



Value-based vs. Policy-based RL

- **Value-based**
 - Learn value function
 - Implicit policy (e.g., ϵ -greedy)
- **Policy-based**
 - No value function
 - Learn policy
- **Actor-critic**
 - Learn value function
 - Learn policy




Advantages of Policy-based RL

- **Advantages:**
 - Good convergence properties
 - Easily extended to high-dimensional or continuous state and action spaces
 - Can learn *stochastic* policies
 - Sometimes policies are simple while values and models are complex
 - e.g., rich domain, but optimal is always to go left
- **Disadvantages:**
 - Susceptible to local optima (especially with non-linear FA)
 - Obtained knowledge is specific, does not always generalize well
 - Ignores a lot of information in the data (when used in isolation)

Stochastic Policies


- We have seen deterministic policies like this:
 - State gives $Q(s, a; w)$ and we selected $\pi(a|s)$ by $\operatorname{argmax}_a Q(s, a; w)$



- Instead, stochastic policies do something like this:

$$\pi(a|s) = \mathbb{P}[a|s; \theta]$$

(policy is represented as a probability distribution)



<https://towardsdatascience.com/self-learning-ai-agents-iv-stochastic-policy-gradients-b53f088fce20>

Why do we need stochastic policies?

Example #1: Rock-Paper-Scissors


- Two-player game of rock-paper-scissors
 - Scissors beats paper
 - Rock beats scissors
 - Paper beats rock
- Consider policies for iterated rock-paper-scissors
 - A deterministic policy (e.g., greedy or even ϵ -greedy) is easily exploited
 - A uniform random policy is the optimal policy (i.e., Nash equilibrium)



David Silver, UCL Lecture on Reinforcement Learning. 2015

Why do we need stochastic policies?

Example #2: Aliased Gridworld




- Consider features of the following form (for all N, E, S, W):

$$\phi(s, a) = \begin{array}{ccccccccc} & \text{walls} & & & \text{actions} & & & & \\ & \overbrace{\quad\quad\quad} & & & \overbrace{\quad\quad\quad} & & & & \\ & 1 & 0 & 1 & 0 & 0 & 1 & 0 & 0 \\ \underbrace{\quad} & \underbrace{\quad} & \underbrace{\quad} & \underbrace{\quad} & \underbrace{\quad} & \underbrace{\quad} & \underbrace{\quad} & \underbrace{\quad} & \underbrace{\quad} \\ & N & E & S & W & N & E & S & W \end{array}$$

- The agent cannot differentiate the grey states
- Compare *deterministic* and *stochastic* policies

Why do we need stochastic policies?


Example #2: Aliased Gridworld



- Value-based RL learns a near-deterministic policy
 - e.g., greedy or ϵ -greedy
- Under aliasing, an optimal *deterministic* policy will either
 - Move W in both grey states (shown by red arrows)
 - Move E in both grey states
- Either way, it can get stuck and never reach the money
- Hence, it will traverse the corridor for a long time

Why do we need stochastic policies?

Example #2: Aliased Gridworld




- Instead,
an optimal *stochastic* policy moves randomly E or W in grey states:

$$\begin{aligned}\pi_{\theta}(\text{wall to N and S, move E}) &= 0.5 \\ \pi_{\theta}(\text{wall to N and S, move W}) &= 0.5\end{aligned}$$

- Will reach the goal state in a few steps with high probability
- Policy-based RL can learn the optimal stochastic policy

Why is it better to learn the policy directly?

- Example: Cartpole




$(x_t, \dot{x}_t, \vartheta, \dot{\vartheta}_t)$



$$a_t = \theta_0 + \theta_1 x_t + \theta_2 \dot{x}_t + \theta_3 \vartheta + \theta_4 \dot{\vartheta}_t$$



$(x_t, \dot{x}_t, \vartheta, \dot{\vartheta}_t)$



Why is it better to learn the policy directly?

- Learn directly a policy without calculating value functions in between
- Why?

- **Greedy updates**

$$\theta_{n+1} = \underset{\theta}{\operatorname{argmax}} E_{\pi_{\theta}} \{Q^{\pi}(s, a)\}$$



Potentially unstable learning process with large policy “jumps”

- **Smooth updates**

$$\theta_{n+1} = \theta_n + \boxed{\alpha_n \nabla G_{\theta_n}}$$



Reminder:

$$G = r_0 + \gamma r_1 + \gamma^2 r_2 + \gamma^3 r_3 + \dots = \sum_{t=0}^{\infty} \gamma^t r_t$$

Stable learning process with smooth policy improvement

Why is it better to learn the policy directly?

- Learn directly a policy without calculating value functions in between
- How to calculate the gradient term?

$$\theta_{n+1} = \theta_n + \alpha_n \nabla G_{\theta_n}$$

- Simple optimization: **Finite Difference Stochastic Approximation (FDSA)**
- Idea: to evaluate the gradient, for each dimension $k \in [1, n]$:
 - Estimate k -th partial derivative of objective function w.r.t. θ by perturbation θ by a small amount ϵ in k -th dimension:

$$\frac{\partial J(\theta)}{\partial \theta_k} \approx \frac{J(\theta + \epsilon u_k) - J(\theta)}{\epsilon},$$

where u_k is a unit vector with 1 in k -th component, 0 elsewhere; $\lim_{n \rightarrow \infty} \epsilon = 0$


- **In RL literature:** “Finite Difference Gradient Estimator”
- **Note:** a variation in control literature is called Simultaneous Perturbation Stochastic Approximation (SPSA)
- Simple, noisy, inefficient – but sometimes effective
 - works for arbitrary policies (even if they are not differentiable)!

Example: AIBO with FDSA

- Learn fast walking patterns for RoboCup (speed decides on win/lose)
- Policy parametrized as an ellipsoid (12 parameters)
- Adapt parameters by (sampled) FDSA
- Policy evaluated by field traversal time



http://www.cs.utexas.edu/users/AustinVilla/?p=research/learned_walk



	π_1	$\pi_2 - \pi_N$	Score	
$-\epsilon_1$	$\theta_1 - \epsilon_1$...	207	\Rightarrow Average: 210
	$\theta_1 - \epsilon_1$...	214	
	...			
$+0$	$\theta_1 + 0$...	225	\Rightarrow Average: 220
	$\theta_1 + 0$...	220	
	...			
$+\epsilon_1$	$\theta_1 + \epsilon_1$...	239	\Rightarrow Average: 240
	$\theta_1 + \epsilon_1$...	244	
	...			

Kohl et al.: Policy gradient reinforcement learning for fast quadrupedal locomotion. ICRA' 2004.

Example: AIBO with FDSA

- Learn fast walking patterns for RoboCup (speed decides on win/lose)
- Policy parametrized as an ellipsoid (12 parameters)
- Adapt parameters by (sampled) FDSA
- Policy evaluated by field traversal time

Problems:

- Requires **A LOT** of samples/trajectories
- In stochastic environments, with small c_n it is really hard to distinguish the difference between R^+ and R^-

Better: Augmented Random Search (ARS)

- Builds on the **Basic Random Search (BRS)** Algorithm:
 - Pick a policy π_θ , perturb the parameters θ by applying $+v\delta$ and $-v\delta$ ($v < 1$ is constant noise and δ is a random number sampled from a normal distribution)
 - Run the policies and apply actions based on $\pi(\theta + v\delta)$ and $\pi(\theta - v\delta)$ and collect the rewards $r(\theta + v\delta)$ and $r(\theta - v\delta)$
 - For all δ compute the average $\Delta = \frac{1}{N} \cdot \Sigma[r(\theta + v\delta) - r(\theta - v\delta)]\delta$ and update the parameters θ using Δ and a learning rate α :

$$\theta_{j+1} = \theta_j + \frac{\alpha}{N} \sum_{k=1}^N [r(\pi_{j,k,+}) - r(\pi_{j,k,-})] \delta_k$$

- Augmented Random Search (ARS) adds 3 improvements:
 - Divide by the rewards by their standard deviation σ_r
 - Normalize the states
 - Only use the top- k best rollouts to compute the average

Better: Augmented Random Search (ARS)

Algorithm 1 Augmented Random Search (ARS): four versions **V1**, **V1-t**, **V2** and **V2-t**

- 1: **Hyperparameters:** step-size α , number of directions sampled per iteration N , standard deviation of the exploration noise ν , number of top-performing directions to use b ($b < N$ is allowed only for **V1-t** and **V2-t**)
- 2: **Initialize:** $M_0 = \mathbf{0} \in \mathbb{R}^{p \times n}$, $\mu_0 = \mathbf{0} \in \mathbb{R}^n$, and $\Sigma_0 = \mathbf{I}_n \in \mathbb{R}^{n \times n}$, $j = 0$.
- 3: **while** ending condition not satisfied **do**
- 4: Sample $\delta_1, \delta_2, \dots, \delta_N$ in $\mathbb{R}^{p \times n}$ with i.i.d. standard normal entries.
- 5: Collect $2N$ rollouts of horizon H and their corresponding rewards using the $2N$ policies

Sample N different variations for the policy parameters (δ_i)

Run $2N$ simulations/rollouts for the positive and negative directions

$$\mathbf{V1}: \begin{cases} \pi_{j,k,+}(x) = (M_j + \nu\delta_k)x \\ \pi_{j,k,-}(x) = (M_j - \nu\delta_k)x \end{cases}$$

$$\mathbf{V2}: \begin{cases} \pi_{j,k,+}(x) = (M_j + \nu\delta_k) \text{diag}(\Sigma_j)^{-1/2} (x - \mu_j) \\ \pi_{j,k,-}(x) = (M_j - \nu\delta_k) \text{diag}(\Sigma_j)^{-1/2} (x - \mu_j) \end{cases}$$

In $\mathbf{Vx-t}$ version of the algorithm, select only the best b rollouts for the parameter update

- for $k \in \{1, 2, \dots, N\}$.
- 6: **V1-t, V2-t:** Sort the directions δ_k by $\max\{r(\pi_{j,k,+}), r(\pi_{j,k,-})\}$, denote by $\delta_{(k)}$ the k -th largest direction, and by $\pi_{j,(k),+}$ and $\pi_{j,(k),-}$ the corresponding policies.
- 7: Make the update step:

To avoid tuning the learning rate, scale the update by the standard deviation (σ_R) of the $2b$ returns used for the update

$$M_{j+1} = M_j + \frac{\alpha}{b\sigma_R} \sum_{k=1}^b [r(\pi_{j,(k),+}) - r(\pi_{j,(k),-})] \delta_{(k)},$$

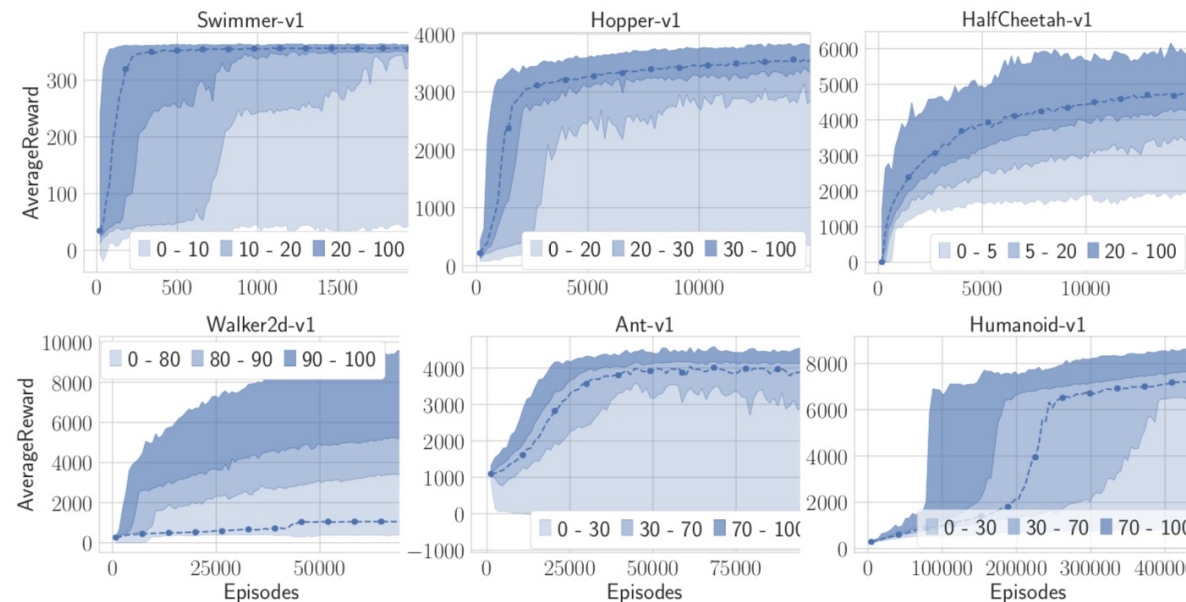
In **V2** of the algorithm, do not use the state observed as input but normalize states using the running mean and variance of all states observed so far

- where σ_R is the standard deviation of the $2b$ rewards used in the update step.
- 8: **V2:** Set μ_{j+1}, Σ_{j+1} to be the mean and covariance of the $2NH(j+1)$ states encountered from the start of training.
- 9: $j \leftarrow j + 1$
- 10: **end while**

Better: Augmented Random Search (ARS)

- State-of-the Art algorithm extending classical random search method
- Comparable performance to modern Deep RL algorithms
- Robust to hyper-parameters and minimum tuning required
- Developed by the Control Engineering Community!

Average reward evaluated over 100 random seeds, shown by percentile




Mania et al.: Simple random search of static linear policies is competitive for reinforcement learning. NeurIPS 2018.

see also: <https://towardsdatascience.com/introduction-to-augmented-random-search-d8d7b55309bd>


Better: Augmented Random Search (ARS)

- State-of-the Art algorithm extending classical random search method
- Comparable performance to modern Deep RL algorithms
- Robust to hyper-parameters and minimum tuning required
- Developed by the Control Engineering Community!


at the beginning...



after 100 iterations:



after 300 iterations:



see also: <https://towardsdatascience.com/introduction-to-augmented-random-search-d8d7b55309bd>

Better: Augmented Random Search (ARS)

Pros:

- Simple to understand and implement
- Less parameter tuning and robust to hyper-parameters
- Embarrassingly easy to parallelize

Cons:

- They tend to favor “lucky” rollouts
- In stochastic environments it is not easy to distinguish if good performance is due to parameter variation or environment noise
- They do not exploit the sequential structure of the problem
- They require 10x more samples (approx.) compared to properly tuned Deep RL algorithms

But...wait...uh... What are we doing here?

