

# Memory-based Exploration

Christopher Mutschler



# Agenda

- Motivation, Problem Definition & Multi-Armed Bandits
- Classic Exploration Strategies
  - Epsilon Greedy
  - (Bayesian) Upper Confidence Bounds
  - Thomson Sampling
- **Exploration in Deep RL**
  - Count-based Exploration: Density Models, Hashing
  - Prediction-based Exploration:
    - Forward Dynamics
    - Random Networks
    - Physical Properties
  - **Memory-based Exploration:**
    - Episodic Memory
    - Direct Exploration
- Summary and Outlook

# Memory-based Exploration

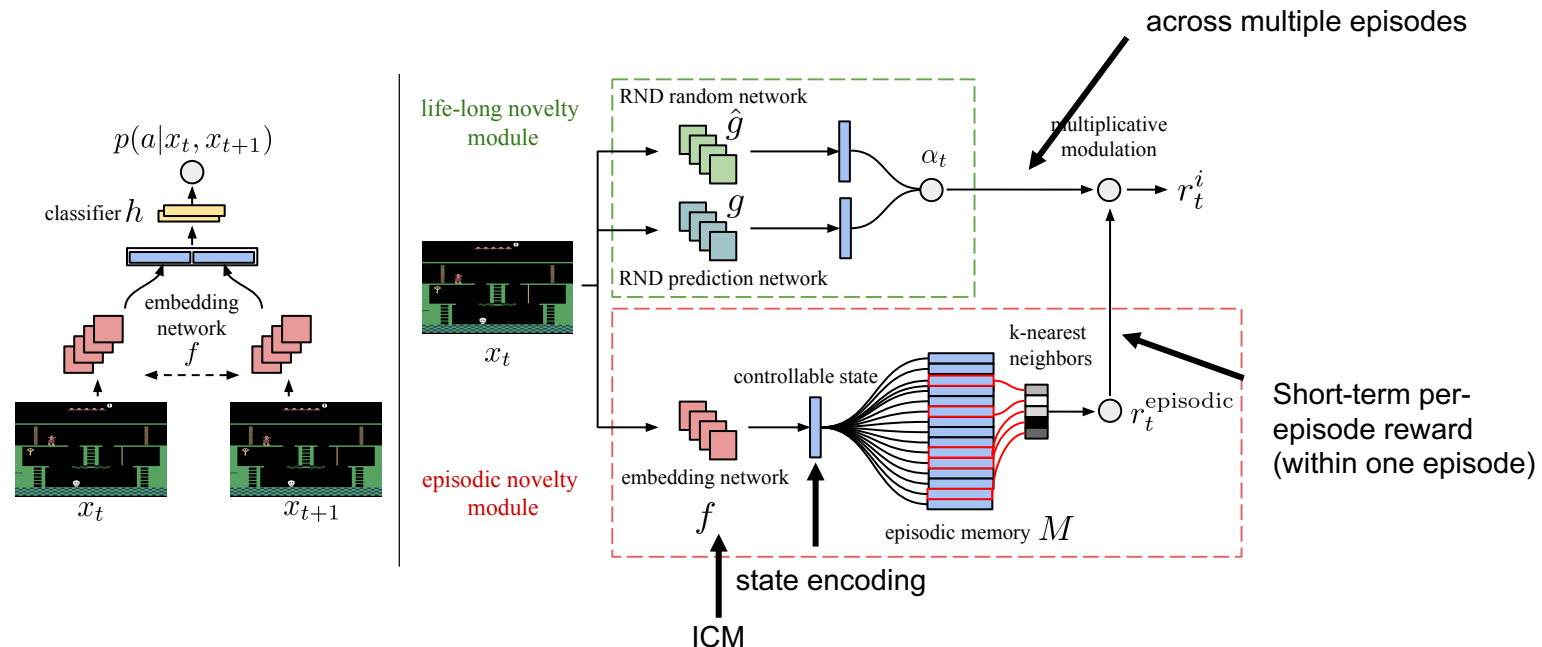
- Reward-based exploration works well in many applications
- However, it suffers from several disadvantages:
  - Function approximation is slow
  - Exploration bonus is non-stationary
  - Knowledge fading: states are no longer novel in time and do no longer provide intrinsic reward signals

## Idea of memory-based exploration:

→ *Use external memories in combination to resolve such disadvantages!*

# Episodic Memory: Never Give Up (NGU)<sup>1</sup>

- So far: RND works great but suffers from episodic settings
- Idea: use two modules:
  1. RND as a lifelong novelty module, and
  2. an episodic novelty module for rapid in-episode adaptation



<sup>1</sup> Badia et al.: Never Give Up: Learning Directed Exploration Strategies. ICLR 2020.

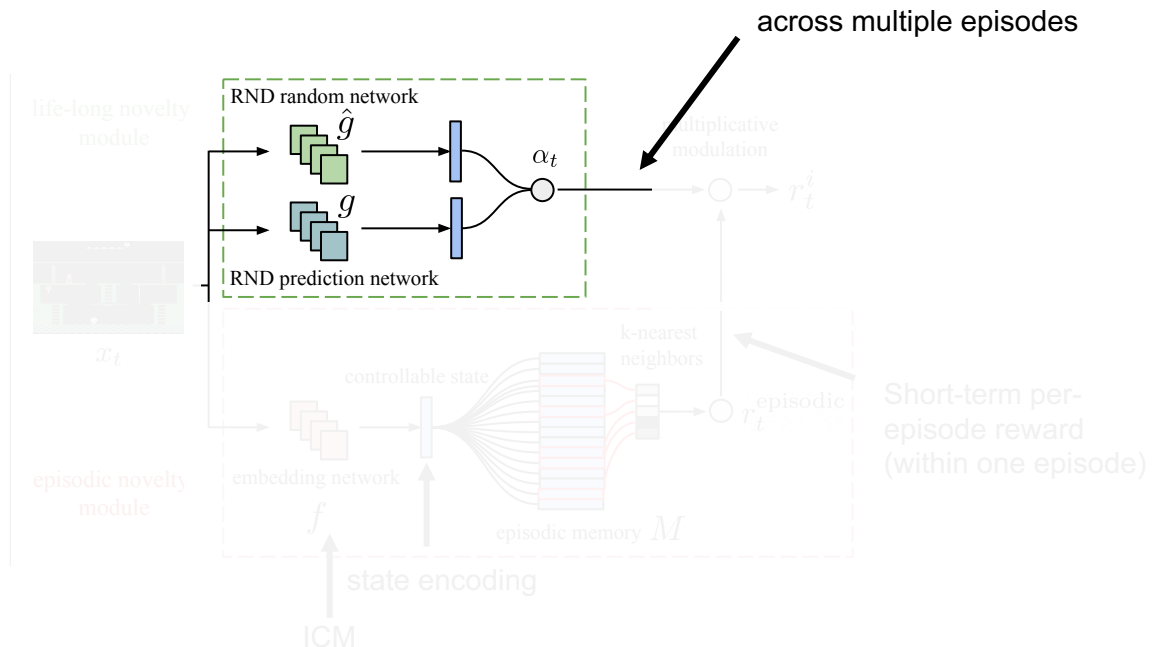
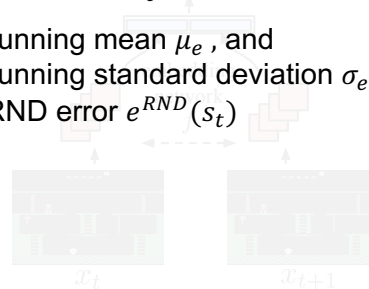
# Episodic Memory: Never Give Up (NGU)<sup>1</sup>

- So far: RND works great but suffers from episodic settings
- Idea: use two modules:
  1. RND as a lifelong novelty module, and
  2. an episodic novelty module for rapid in-episode adaptation

- RND derives a life-long novelty bonus
- The exploration bonus is given as

$$\alpha_t = 1 + \frac{e^{RND(s_t)} - \mu_e}{\sigma_e}, \text{ with}$$

- the running mean  $\mu_e$ , and
- the running standard deviation  $\sigma_e$  of the RND error  $e^{RND}(s_t)$



<sup>1</sup> Badia et al.: Never Give Up: Learning Directed Exploration Strategies. ICLR 2020.

# Episodic Memory: Never Give Up (NGU)<sup>1</sup>

- So far: RND works great but suffers from episodic settings
- Idea: use two modules:
  1. RND as a lifelong novelty module, and
  2. an episodic novelty module for rapid in-episode adaptation

- Add  $\phi(s_t)$  into  $M$
- Compare  $\phi(s_t)$  to the other content in  $M$ :


$$r_t^e = \frac{1}{\sqrt{\sum_{\phi_i \in N_k} K(\phi(x_t), \phi_i) + c}}, \text{ with}$$

- a kernel function  $K(x, y)$
- # nearest neighbors  $N_k$ , and
- a constant  $c$  to avoid a zero-sum

Originally, the paper proposes

$$K(x, y) = \frac{\epsilon}{\frac{d^2(x, y)}{d_m^2} + \epsilon}, \text{ with}$$

- Euclidean distance  $d$  between two samples
- squared Euclidean distance  $d_m^2$  of the  $k$ -th nearest neighbors  $\rightarrow$  more robust
- a small constant  $\epsilon$



<sup>1</sup> Badia et al.: Never Give Up: Learning Directed Exploration Strategies. ICLR 2020.

# Episodic Memory: Never Give Up (NGU)<sup>1</sup>

- So far: RND works great but suffers from episodic settings
- Idea: use two modules:
  1. RND as a lifelong novelty module, and
  2. an episodic novelty module for rapid in-episode adaptation

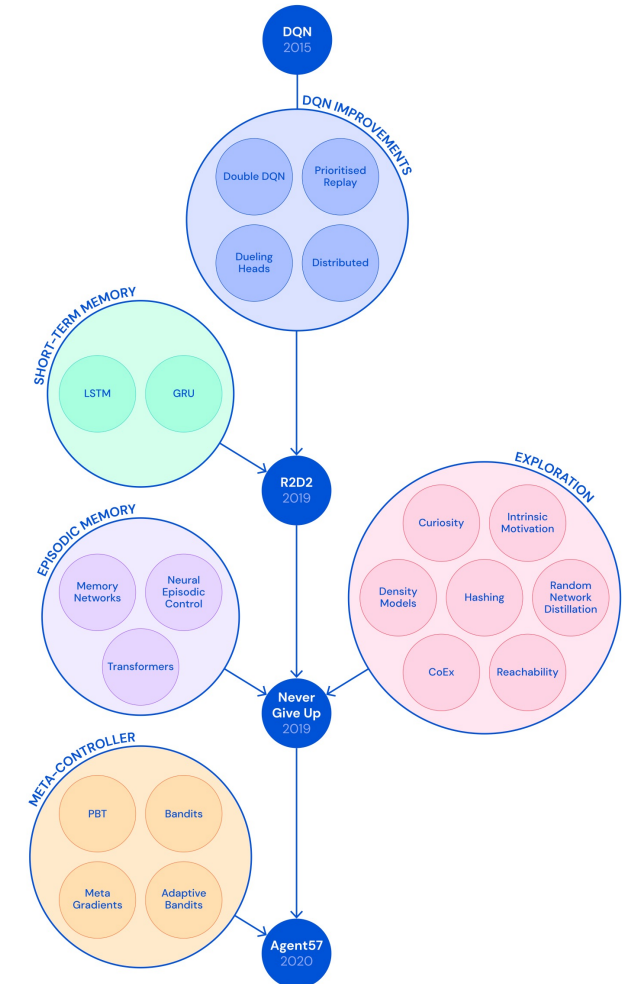


→ **Do not revisit the same state within the same episode!**  
 → **Try to not revisit the states you already saw in previous episodes!**

<sup>1</sup> Badia et al.: Never Give Up: Learning Directed Exploration Strategies. ICLR 2020.

# Agent57<sup>1</sup>

- Agent57 is the first RL agent who beats Atari57 consistently
- Two main improvements over NGU:
  1. Population of policies:
    - Each policy has its own pair of exploration parameters  $\{(\beta_j, \gamma_j)\}_{j=1}^N$
    - Policies with high  $\beta_j$  (and lower  $\gamma_j$ ) make more progress at early stages
    - Policies with high  $\gamma_j$  (and lower  $\beta_j$ ) make more progress at later stages
    - A meta-controller (sliding window UCB) is trained to select from the policies
  2. Re-Parameterization of Q-value function:
    - Q-function is decomposed into intrinsic and extrinsic influence:
 
$$Q(s, a; \theta_j^i) = Q(s, a; \theta_j^e) + \beta_j Q(s, a; \theta_j^i)$$
    - During training both parameter sets ( $\theta^e$  and  $\theta^i$ ) are optimized separately with rewards  $r_j^e$  and  $r_j^i$ , respectively




<https://deepmind.com/blog/article/Agent57-Outperforming-the-human-Atari-benchmark>

<sup>1</sup> Badia et al.: Agent 57: Outperforming the Atari Human Benchmark. ICML 2020.



# Episodic Curiosity through Reachability<sup>1</sup>



- There is a better thing than using the Euclidean distance\*
- Episodic Curiosity (EC) module:
  - Measure the number of steps needed to transit between two states
  - The novelty then depends on the *reachability* between states
- General idea/steps:
  1. Clear episodic memory  $M$  on environment reset
  2. At each step  $t$  until the episode ends:
    1. Compare  $s_t$  with all the states in the memory
    2. If it takes more than  $k$  steps to reach  $s_t \rightarrow$  agent gets a bonus
    3. If the novelty bonus is large enough  $\rightarrow$  add  $s_t$  to  $M$
- But how can we estimate the *reachability*?



<sup>1</sup> Savinov et al.: Episodic Curiosity through Reachability. ICLR 2019.

# Episodic Curiosity through Reachability<sup>1</sup>

- Ideally, we would have access to a transition graph
  1. Not possible to build up (due to limited memory)
  2. Hard to build
- Solution: Train a Siamese neural network that predicts how far two states are apart
  - Embedding network  $E: \mathcal{O} \rightarrow \mathbb{R}^n$   
(encodes states to feature vectors)
  - Comparator network  $C: \mathbb{R}^n \times \mathbb{R}^n \rightarrow [0,1]$   
(reachability within  $k$  steps: 0 (not reachable) to 1 (reachable))
  - “R-network” as a classifier trained with logistic regression loss, based on trajectory data:  
→ Hence:  $R(o_i, o_j) = C(E(o_i), E(o_j))$



<sup>1</sup> Savinov et al.: Episodic Curiosity through Reachability. ICLR 2019.

# Episodic Curiosity through Reachability

- Putting all together: **Episodic Curiosity (EC) Module**

- At every time step

1. Embedding network processes  $o_t \rightarrow$  embedding vector  $e = E(o_t)$

2. Compare  $e$  with all embeddings in the buffer  $M = \langle e_1, \dots, e_{|M|} \rangle$  via  $C \rightarrow$  fills the reachability buffer with values

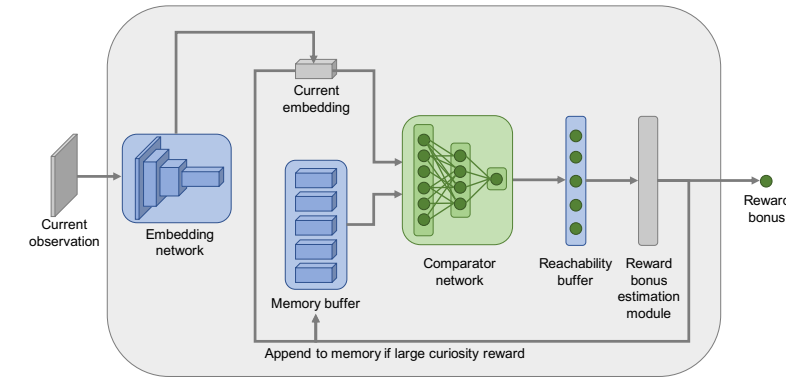
$$c_i = C(e_i, e), \quad i = 1, \dots, M.$$

3. Compute the similarity score between  $e$  and the memory buffer  $M$  as  $C(M, e) = F(c_1, \dots, c_{|M|}) \in [0, 1]$ , where  $F(\cdot)$  is a hyperparameter (function).

- $\max(\cdot)$  would theoretically be a good choice but is prone to outliers
- 90<sup>th</sup> percentile works better in experiments.

4. Compute the curiosity bonus as  $b = B(M, e) = \alpha(\beta - C(M, e))$

- $\alpha$  tunes scale of task rewards
- $\beta$  defines the sign of the reward (0.5 works well for fixed-duration episodes)




# Direct Exploration

## Go-Explore<sup>1</sup>

- The problems stemming from sparse rewards and intrinsic motivation are two-fold:
  - Detachment**
    - Intrinsic rewards are nearly always a consumable resource: short-term focus lies on such areas but with time they become less interesting to the agent
    - Catastrophic forgetting: we forget things that happened far in the past*
  - Derailment**
    - Describes the problem of re-visiting an interesting state again, in order to further explore from there
    - Previous work runs the policy again (with stochastic perturbation) in a "hope" to reach the desired state again

→ with naïve perturbations in complex environments this becomes highly unlikely



see also: <https://towardsdatascience.com/a-short-introduction-to-go-explore-c61c2ef201f0>

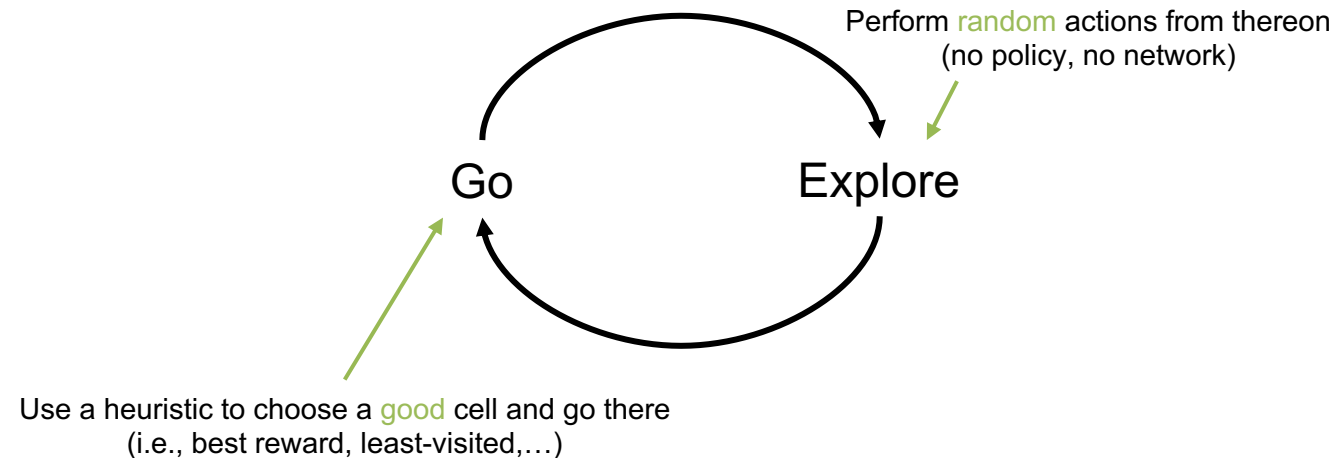


<sup>1</sup> Adrian Ecoffet et al.: Go-Explore: a new Approach for Hard-Exploration Problems. 2020.

# Direct Exploration: Go-Explore

**Go-Explore**<sup>1</sup> addresses *detachment* and *derailment* using two phases:


## 1. Explore until solved

- No ML and NNs involved here. Just random (or semi-guided) exploration
- Main goal: find **interesting cells**
  - Newly discovered & high reward obtained to reach them
- For each interesting cell we store the
  1. full trajectory to get there
  2. a snapshot of the environment state
  3. total reward of the trajectory
  4. length of the trajectory
- If we revisit a state
  - Update the entry if it is better (i.e., short trajectory, higher reward)



<sup>1</sup> Adrian Ecoffet et al.: Go-Explore: a new Approach for Hard-Exploration Problems. 2020.

# Direct Exploration: Go-Explore



**Go-Explore**<sup>1</sup> addresses *detachment* and *derailment* using two phases:

- Limitation of 1<sup>st</sup> phase: go to a cell is only “easy” in deterministic environments!

## 2. Robustify (if needed): “Backward” algorithm

- Consider a sequence  $c_1, c_2, \dots, c_{n-1}, c_n$ , where  $c_n$  is the “go” cell
- Algorithm:
  1. Initialize  $c_i = c_{n-1}$
  2. Set environment to the snapshot of  $c_i$  and train the agent to reach  $c_n$
  3. When agent finds a trajectory with an equal or higher reward:
    - set  $i \leftarrow i - 1$
    - go to step 2
  4. Stop when  $i = 1$
- How to make policies more robust to non-determinism in Atari?
  - No ops
  - Sticky actions

<sup>1</sup> Adrian Ecoffet et al.: Go-Explore: a new Approach for Hard-Exploration Problems. 2020.

# Direct Exploration: Go-Explore


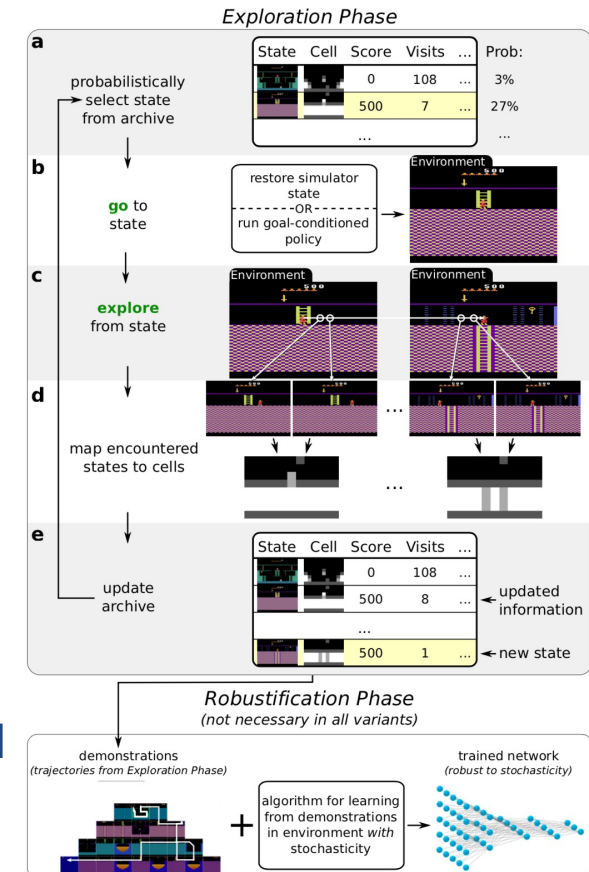


Figure 6: **History of progress on Montezuma's Revenge vs. the version of Go-Explore that does not harness domain knowledge.** Go-Explore significantly improves on the prior state of the art.

<sup>1</sup> Adrian Ecoffet et al.: Go-Explore: a new Approach for Hard-Exploration Problems. 2020.

# Direct Exploration: Go-Explore (Improvements)

- “First return, then explore”<sup>1</sup>: *policy-based Go-Explore*
  - Instead of resetting the simulator: learn a goal-conditioned policy to reach a state  
→ mainly trained to follow the best trajectory so far
  - Self-Imitation Learning to extract more information from successful trajectories
- “Memory based Trajectory-conditioned Policies for Learning from Sparse Reward”
  - Like policy-based go-explore:
    - Maintain a memory of demonstrations collected during training
    - Use them to train a trajectory-conditioned policy via Self-Imitation Learning
    - Prioritize trajectories that end with a rare state during sampling.



<sup>1</sup> Adrian Ecoffet et al.: First return, then explore. 2021.

<sup>2</sup> Yijie Guo et al.: Memory Based Trajectory-conditioned Policies for Learning from Sparse Rewards. 2021.



# Agenda

- Motivation, Problem Definition & Multi-Armed Bandits
- Classic Exploration Strategies
  - Epsilon Greedy
  - (Bayesian) Upper Confidence Bounds
  - Thomson Sampling
- Exploration in Deep RL
  - Count-based Exploration: Density Models, Hashing
  - Prediction-based Exploration:
    - Forward Dynamics
    - Random Networks
    - Physical Properties
  - Memory-based Exploration:
    - Episodic Memory
    - Direct Exploration
- **Summary and Outlook**

# Summary

- We studied different families of algorithms and settings in this lecture:
  1. Multi-armed Bandits and their theoretical assumptions
  2. Challenges that arise from going from small MDPs to high-dimensional POMDPs
  3. We found different families of methods to guide exploration in Deep RL:
    - Count-based Exploration
    - Prediction-based Exploration
    - Memory-based Exploration
- Exploration is **really hot topic** in current RL research
  - Proving theoretical assumption and bound from (contextual) bandits on small MDPs, as well as
  - Exploration in Deep RL

# References

- Lilian Weng: The Multi-Armed Bandit Problem and Its Solutions. lilianweng.github.io/lil-log, 2018.  
<https://lilianweng.github.io/lil-log/2018/01/23/the-multi-armed-bandit-problem-and-its-solutions.html>
- CS229 Supplemental Lecture notes. <http://cs229.stanford.edu/extra-notes/hoeffding.pdf>
- UCL Course by David Silver – Lecture 9: XX.
- Oliver Chapelle and Lihong Li: An empirical evaluation of Thompson Sampling. NIPS2011.
- Daniel Russo et al.: A Tutorial on Thompson Sampling. arXiv:1707.02038. 2017.
- <https://openai.com/blog/reinforcement-learning-with-prediction-based-rewards/>