

Generative Adversarial Networks

Advanced Topics in Machine Learning

Andrea Cini, Cesare Alippi

22 September, 2021

Università della Svizzera italiana
slides credits: Daniele Grattarola

Which photo is real?



Which photo is real?



thispersondoesnotexist.com

Generative Adversarial Networks [1] in short

Goal: generate samples that **look like** the real data.

[1] I. J. Goodfellow et al., *Generative Adversarial Networks*, 2014.

Generative Adversarial Networks [1] in short

Goal: generate samples that **look like** the real data.

Two main **neural networks**:

- **Generator** $G(z)$ maps random noise z to the data space;
- **Discriminator** $D(x)$ decides whether sample x was generated by G or is a real sample;

[1] I. J. Goodfellow et al., *Generative Adversarial Networks*, 2014.

Generative Adversarial Networks [1] in short

Goal: generate samples that **look like** the real data.

Two main **neural networks**:

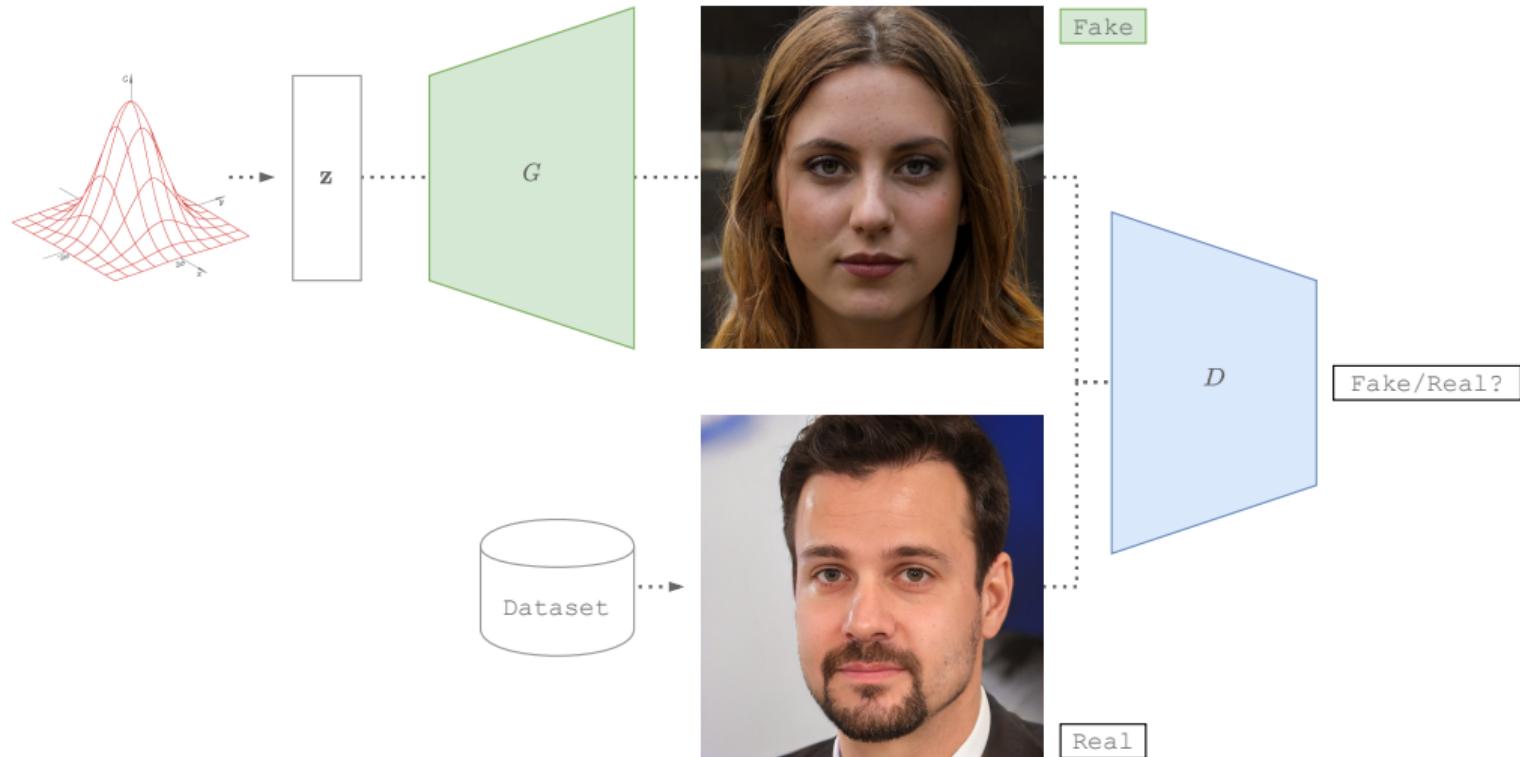
- **Generator** $G(z)$ maps random noise z to the data space;
- **Discriminator** $D(x)$ decides whether sample x was generated by G or is a real sample;

The two components play **against each other** until the **generator** fools the **discriminator**.

The diagram shows the text "The two components play **against each other** until the **generator** fools the **discriminator**." with a brace underneath the words "against each other" and "generator". Below the brace, the word "Adversarial" is written in a smaller font.

[1] I. J. Goodfellow et al., *Generative Adversarial Networks*, 2014.

Generative Adversarial Networks in short



Different types of neural networks can be used as **generator**:

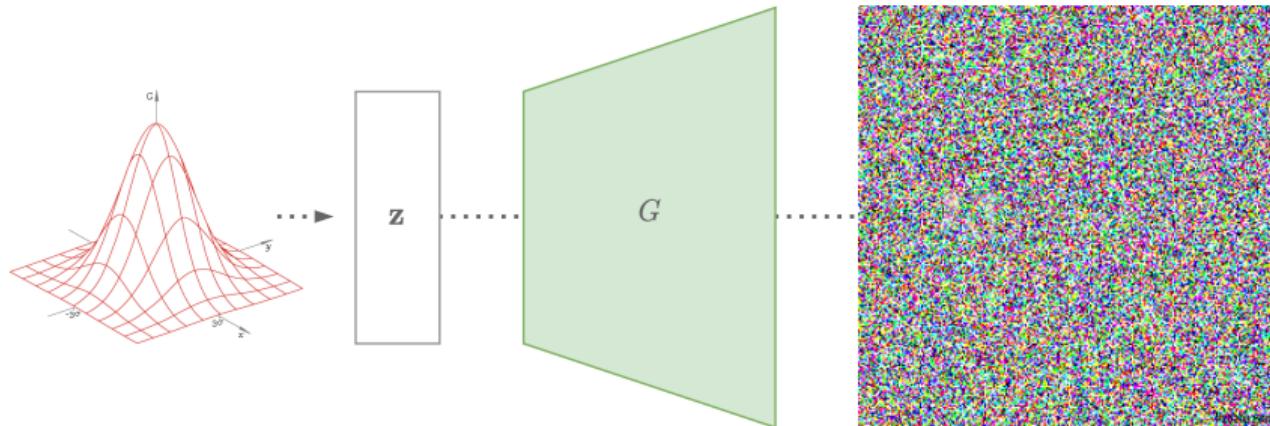
- Convolutional for images or audio;
- Recurrent for text or sequences;
- MLPs for tabular data;

Different types of neural networks can be used as **generator**:

- **Convolutional for images or audio;** ← we focus mostly on these (easier to visualize)
- Recurrent for text or sequences;
- MLPs for tabular data;

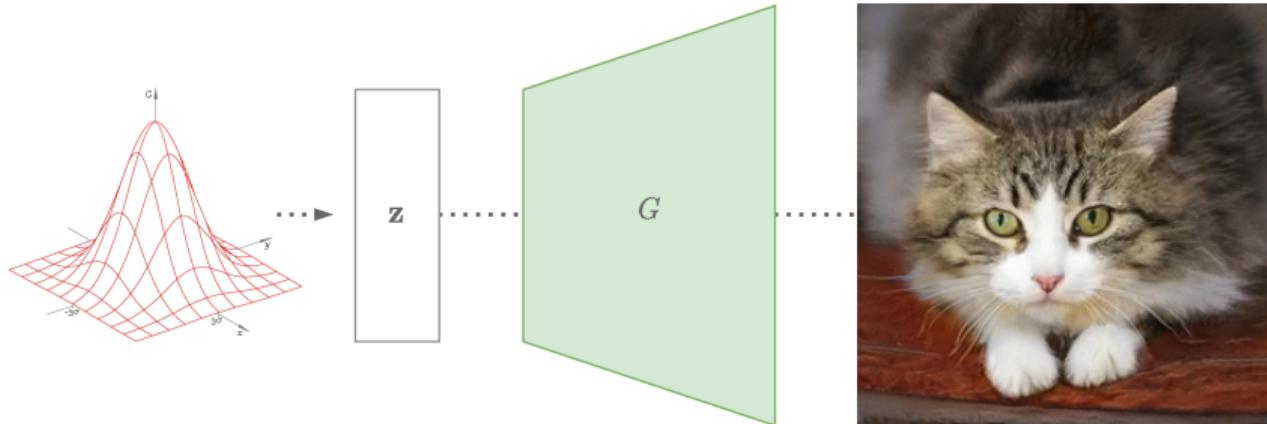
Generator

Generates an image of w by h pixels $\mathbf{x} = G(\mathbf{z}) \in [0, 1]^{w \times h}$, where $\mathbf{z} \sim p_z(\mathbf{z})$ (e.g., $\mathcal{N}(0, I)$).



Generator

Generates an image of w by h pixels $\mathbf{x} = G(\mathbf{z}) \in [0, 1]^{w \times h}$, where $\mathbf{z} \sim p_z(\mathbf{z})$ (e.g., $\mathcal{N}(0, I)$).



(thiscatdoesnotexist.com)

Discriminator

The **discriminator** looks at samples from the data space $x \in [0, 1]^{w \times h}$ and outputs:

- 1 if the samples come from the real data distribution, *i.e.*, $x \sim p^*(x)$;
- 0 if the samples are **fake**, *i.e.*, $x \sim p_g(x)$ (*first* $z \sim p_z(z)$, *then* $x = G(z)$);

Discriminator

The **discriminator** is trained to optimise two objectives:

- $\max_D \mathbb{E}_{x \sim p^*(x)} [\log D(x)]$ (output 1 on real samples)
- $\max_D \mathbb{E}_{z \sim p_z(z)} [\log (1 - D(G(z)))]$ (output 0 on **generated** samples)

(this is just a different way to write a binary classification problem)

Training the generator

Recall: the **generator** has to fool the **discriminator**.

$$\underbrace{\max_D \mathbb{E}_{z \sim p_z(z)} [\log (1 - D(G(z)))]}_{\text{Discriminator objective}}$$

Training the generator

Recall: the **generator** has to fool the **discriminator**.

$$\underbrace{\max_D \mathbb{E}_{z \sim p_z(z)} [\log (1 - D(G(z)))]}_{\text{Discriminator objective}} \rightarrow \underbrace{\min_G \mathbb{E}_{z \sim p_z(z)} [\log (1 - D(G(z)))]}_{\text{Generator objective}}$$

Fooling = making the **discriminator** output 1 on generated samples.

Putting everything together

If we combine the two objectives for $G(z)$ and $D(x)$ we get **the GAN min-max game**:

$$\min_G \max_D \mathbb{E}_{x \sim p^*(x)} [\log D(x)] + \mathbb{E}_{z \sim p_z(z)} [\log (1 - D(G(z)))]$$

Training the GAN

The GAN is trained by an iterative procedure, **repeated to convergence**:

1. Train the **discriminator** on a batch of real and fake samples;
2. Train the **generator** to fool the discriminator.

Training the GAN - Discriminator training

For k steps do:

1. Sample a minibatch of noise samples $\{z^{(1)}, \dots, z^{(m)}\}$ from $p_z(z)$;
2. Sample a minibatch of real samples from the dataset $\{x^{(1)}, \dots, x^{(m)}\}$;
3. Update the weights θ_D of D by gradient **ascent**:¹

$$\nabla_{\theta_D} \frac{1}{m} \sum_{i=1}^m \left[\log D(x^{(i)}) + \log \left((1 - D(G(z^{(i)})) \right) \right].$$

¹In practice we can also do gradient descent by changing the sign.

Training the GAN - Generator training

Do once:

1. Sample a minibatch of noise samples $\{z^{(1)}, \dots, z^{(m)}\}$ from $p_z(z)$;
2. Update the weights θ_G of G by gradient **descent**:

$$\nabla_{\theta_G} \frac{1}{m} \sum_{i=1}^m \log \left((1 - D(G(z^{(i)})) \right).$$

Convergence

The min-max game have a single global minimum, achieved *iff* G perfectly fools D .

Convergence

The min-max game have a single global minimum, achieved iff G perfectly fools D .

When this happens:

- The data distribution is **indistinguishable** from the distribution of the generated data:
 $p^* = p_g$.
- The global minimum of the training criterion is:

$$\mathbb{E}_{x \sim p^*(x)} \left[\underbrace{\log D(x)}_{\log(0.5)} \right] + \mathbb{E}_{z \sim p_z(z)} \left[\underbrace{\log (1 - D(G(z)))}_{\log(0.5)} \right] = -\log(4)$$

Convergence

The min-max game have a single global minimum, achieved iff G perfectly fools D .

When this happens:

- The data distribution is **indistinguishable** from the distribution of the generated data:
 $p^* = p_g$.
- The global minimum of the training criterion is:

$$\mathbb{E}_{x \sim p^*(x)} \left[\underbrace{\log D(x)}_{\log(0.5)} \right] + \mathbb{E}_{z \sim p_z(z)} \left[\underbrace{\log (1 - D(G(z)))}_{\log(0.5)} \right] = -\log(4)$$

i.e., the **discriminator** is **maximally confused** (can only output 0.5).

Convergence

The min-max game have a single global minimum, achieved iff G perfectly fools D .

When this happens:

- The data distribution is **indistinguishable** from the distribution of the generated data:
 $p^* = p_g$.
- The global minimum of the training criterion is:

$$\mathbb{E}_{x \sim p^*(x)} \left[\underbrace{\log D(x)}_{\log(0.5)} \right] + \mathbb{E}_{z \sim p_z(z)} \left[\underbrace{\log (1 - D(G(z)))}_{\log(0.5)} \right] = -\log(4)$$

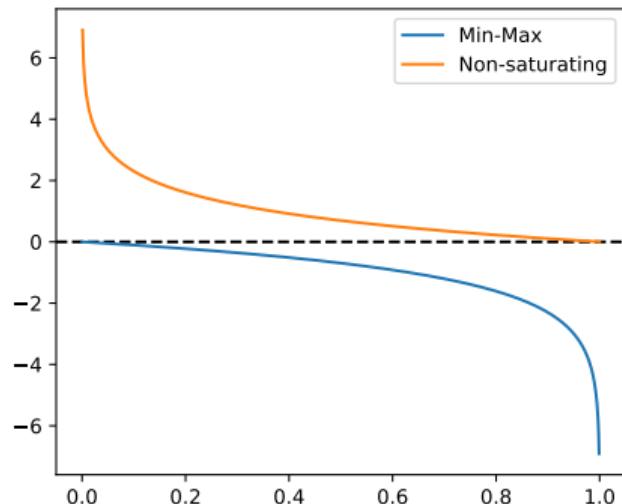
i.e., the **discriminator** is **maximally confused** (can only output 0.5).

Note: convergence is guaranteed only in function space and in the limit of data, in practice we are limited by the amount of data available and the parametrization of G and D .

Tips and Tricks

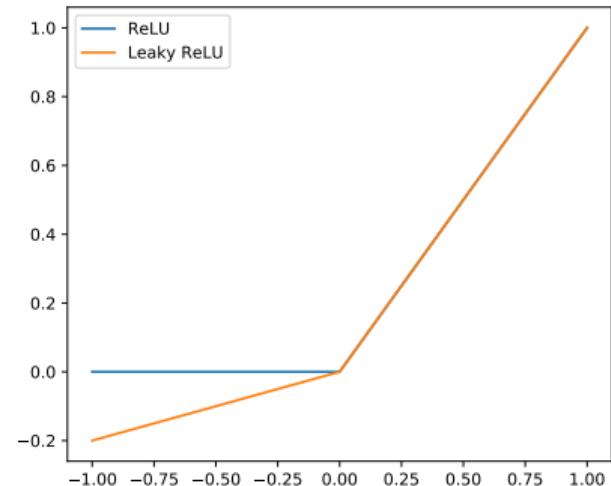
- Use **non-saturating** loss to optimize G :

$$\nabla_{\theta_G} \frac{1}{m} \sum_{i=1}^m -\log D(G(z^{(i)}));$$



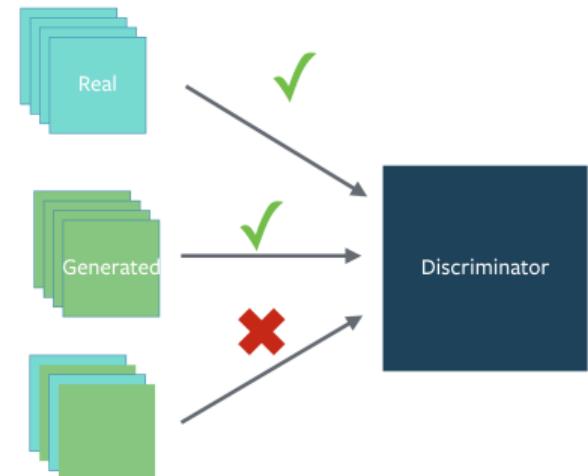
Tips and Tricks

- Use **non-saturating** loss to optimize G :
$$\nabla_{\theta_G} \frac{1}{m} \sum_{i=1}^m -\log D(G(z^{(i)}));$$
- Avoid **sparse gradients**: use LeakyReLU, average pooling.



Tips and Tricks

- Use **non-saturating** loss to optimize G :
$$\nabla_{\theta_G} \frac{1}{m} \sum_{i=1}^m -\log D(G(z^{(i)}));$$
- Avoid **sparse gradients**: use LeakyReLU, average pooling.
- **Do not mix** real and fake samples in the same mini-batch to train D ;



Tips and Tricks

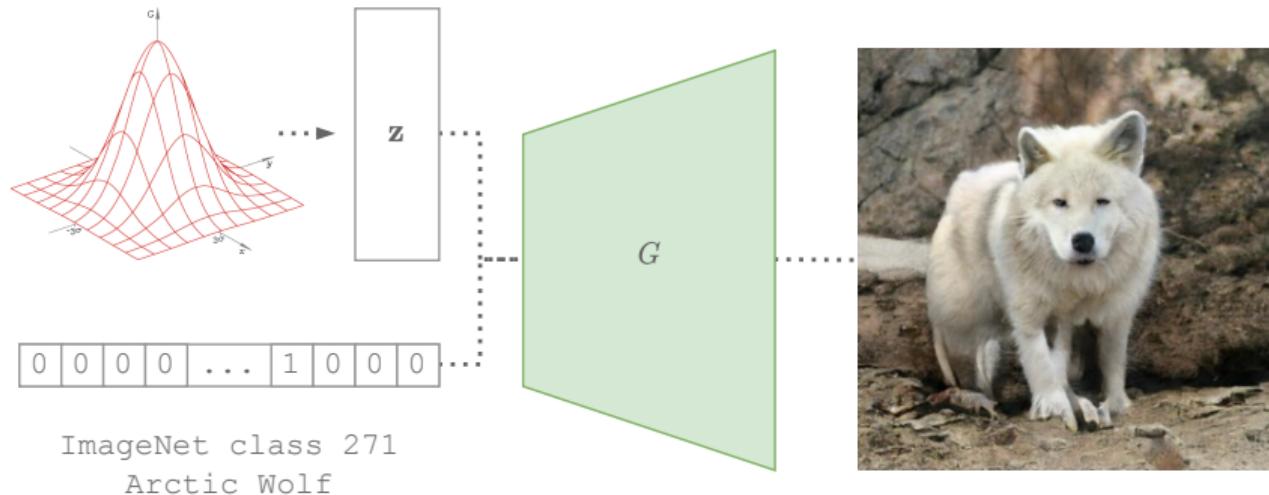
- Use **non-saturating** loss to optimize G :
$$\nabla_{\theta_G} \frac{1}{m} \sum_{i=1}^m -\log D(G(z^{(i)}));$$
- Avoid **sparse gradients**: use LeakyReLU, average pooling.
- **Do not mix** real and fake samples in the same mini-batch to train D ;
- Check out github.com/soumith/ganhacks for more ~~dark~~magic empirical tips.



Conditional GANs

In many cases, the samples from p^* are divided in classes (e.g., ImageNet).

Instead of generating **any** image from p^* , we generate $x = G(z, y)$, where y is a **class label**.



Applications

Text to image translation [2]

The small bird has a red head with feathers that fade from red to gray from head to tail



This bird is black with green and has a very short beak



[2] S. Reed et al., "Generative adversarial text to image synthesis," 2016.

Style transfer [3]



[3] G. Antipov et al., "Face aging with conditional generative adversarial networks," 2017.

Super-resolution [4]



[4] C. Ledig et al., "Photo-realistic single image super-resolution using a generative adversarial network," 2016.

Demo

Demo



cutt.ly/9Eaqevz

Questions?

References i

- [1] I. J. Goodfellow, J. Pouget-Abadie, M. Mirza, B. Xu, D. Warde-Farley, S. Ozair, A. Courville, and Y. Bengio, *Generative adversarial networks*, 2014. arXiv: 1406.2661 [stat.ML].
- [2] S. Reed, Z. Akata, X. Yan, L. Logeswaran, B. Schiele, and H. Lee, “Generative adversarial text to image synthesis,” *arXiv preprint arXiv:1605.05396*, 2016.
- [3] G. Antipov, M. Baccouche, and J.-L. Dugelay, “Face aging with conditional generative adversarial networks,” in *2017 IEEE international conference on image processing (ICIP)*, IEEE, 2017, pp. 2089–2093.
- [4] C. Ledig, L. Theis, F. Huszár, J. Caballero, A. Cunningham, A. Acosta, A. Aitken, A. Tejani, J. Totz, and Z. Wang, “Photo-realistic single image super-resolution using a generative adversarial network,” *arXiv preprint arXiv:1609.04802*, 2016.