

 Università della Svizzera italiana	Institute of Computing CI

High-Performance Computing

2021

Student: Ashutosh Singh

Discussed with: Self

Solution for Project 5

Due date: 21.11.2021, 23:59

HPC 2021 — Submission Instructions

(Please, notice that following instructions are mandatory:
submissions that don't comply with, won't be considered)

- Assignments must be submitted to **iCorsi** (i.e. in electronic format).
- Provide both executable package and sources (e.g. C/C++ files, Matlab). If you are using libraries, please add them in the file. Sources must be organized in directories called:
Project_number_lastname_firstname
 and the file must be called:
project_number_lastname_firstname.zip
project_number_lastname_firstname.pdf
- The TAs will grade your project by reviewing your project write-up, and looking at the implementation you attempted, and benchmarking your code's performance.
- You are allowed to discuss all questions with anyone you like; however: (i) your submission must list anyone you discussed problems with and (ii) you must write up your submission independently.

1. Ring maximum using MPI [10 Points]

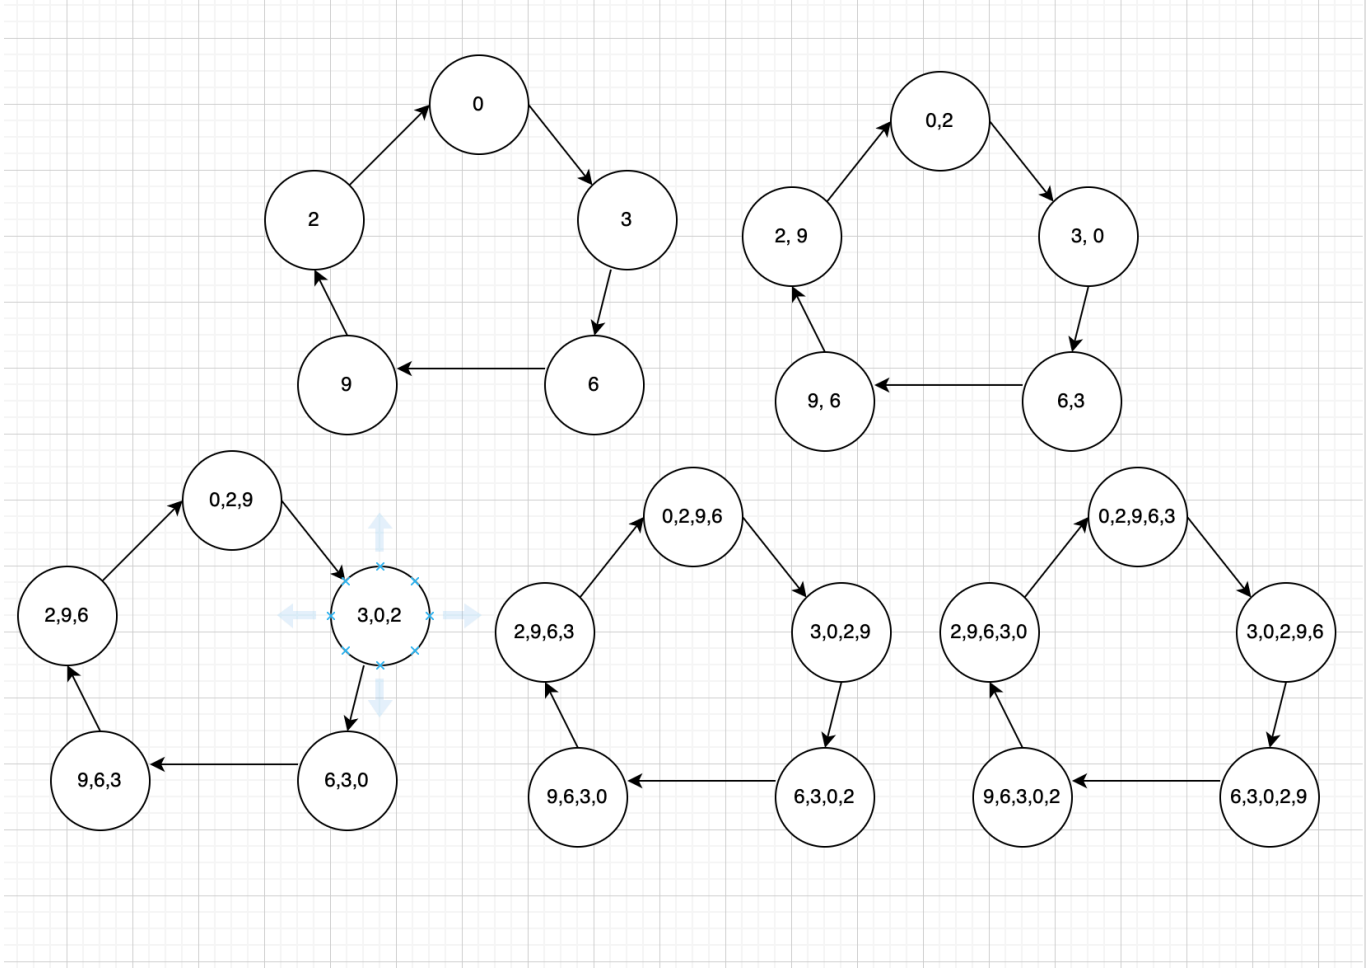


Figure 1: Grid Sizes Vs Log of Runtime(seconds) for Various number of threads. [Click here for interactive figure](#)

As clear from the **Figure-1** each process has two neighbors, one of the left and one on right. The ranks of the neighbours processes are determined by considering a $period = Communicator_Size$. The ranks of right and left neighbours of a process with $rank = my_rank$ are given as -

$$right = my_rank + 1 \bmod Communicator_Size$$

$$left = my_rank + Communicator_Size - 1 \bmod Communicator_Size.$$

At the start of the operation each $process_i$ has its send buffer filled according to the formula. For reference let's call it the original buffer of this process.

$$3 * process_rank_i \bmod 2 * communicator_size, \text{ for } i = 0, 1, \dots, communicator_size - 1.$$

After the first the iteration the send buffer of each process is copied to receive buffer of the right neighbour. After $communicator_size - 1$ such iterations each process has once seen the original send buffer of all other processes.

I have used `MPI_Isend` to send the data from one process to another to prevent the deadlock. Each process has a *result* variable where it stores the max of value in *result* and the receive buffer.

$$result = \max\{result, receive_buffer\}$$

When all the iterations are complete every process has the same value of *result*.

2. Ghost cells exchange between neighboring processes [15 Points]

The computation in this problem is defined on a 2D Cartesian Grid. The processes have to act as if they are 2-Dimensional ordered and they need to communicate with their north, south, east and west neighbours.

The mechanism to declare the structure of a computation to MPI is known as a virtual topology and for this problem the structure is a Cartesian Topology where a process has two neighbours in each dimension.

In Cartesian topology each process is defined by a co-ordinate on the given Cartesian plane. The Cartesian plane itself is described as a space with d dimensions and having $\{n_1, n_2, \dots, n_d\}$ length along each dimension.

In MPI `MPI_Cart_create` method is used to create a Cartesian topology.

```
1 int MPI_Cart_create(MPI_Comm comm_old,
2                     int ndims, int *dims, int *periods,
3                     int reorder, MPI_Comm *comm_cart)
```

For this problem the 2D Cartesian topology we create is *periodic* in both dimensions which means that the rows and columns at the edges round around themselves. So first row has last row as its neighbour and similarly for the first and last column. Also we do not re-order the ranks of the processes.

For each process to find the rank of its N/S/E/W neighbours we use `MPI_Cart_shift`. Snippet from my solution -

```
1 MPI_Cart_shift(comm_cart, 1, +1, &rank_left, &rank_right);
2 MPI_Cart_shift(comm_cart, 0, -1, &rank_bottom, &rank_top);
```

Now to send the data from one process to another I use `MPI_Isend` and `MPI_Irecv`




Figure 2: The Cartesian topology in problem(left) and one process of the topology(right) , each circular element(blue, green and orange) represents one block of memory or one element in a 2D array stored in a row major format.

As can be seen from the [Figure-2](#) each process works on array with $DOMAINSIZE * DOMAINSIZE$ elements. This array is stored in a row major format instead of a 2D array. While communicating with other processes, *SUBDOMAIN* number of elements from top(in blue) are to be sent to the N-neighbour and similarly the bottom elements to the S-neighbour; sending to N/S neighbours is

straightforward because a whole row is to be transmitted and the data array is stored in row major format so it is contiguous for each row. Snippet from code -

```
1 MPI_Isend(&data[1], SUBDOMAIN, MPI_DOUBLE, rank_top, ping, comm_cart, &request);
2 MPI_Isend(&data[DOMAINSIZE * DOMAINSIZE - 1 - SUBDOMAIN], SUBDOMAIN, MPI_DOUBLE,
  rank_bottom, ping, comm_cart, &request);
```

However for W/E neighbours it is not so simple since the blue points(from [Figure](#)) on the left and right and not stored as contiguous block of memory, but the data needs to be accessed one blue point at a time after a stride of *DOMAINSIZE* elements. MPI provides a method to deal with such patterns, called *MPI_Type_vector*. Code Snippet depicting data access and transmission to E-Neighbour.

```
1 MPI_Type_vector(SUBDOMAIN, 1, DOMAINSIZE, MPI_DOUBLE, &data_ghost);
2 MPI_Type_commit(&data_ghost);
3 MPI_Isend(&data, 1, data_ghost, rank_left, ping, comm_cart, &request);
```

In the first line, *MPI_Type_vector* reads 1 memory block(blue element) at a time with a stride of *DOMAINSIZE*, *SUBDOMAIN* number of times from the *data* array (this array represents data in a row major format in one process). In the same way data is transmitted to the W-Neighbour.

The procedure is done for all processes - all sixteen blocks from the picture of left in [Figure-2](#).

Code for receiving data is fairly straightforward and submitted with the solution. For E/W neighbours, they receive data in an array of size *SUBDOMAIN* i.e. the blocklength used by *MPI_Type_vector*.

3. Parallelizing the Mandelbrot set using MPI [20 Points]

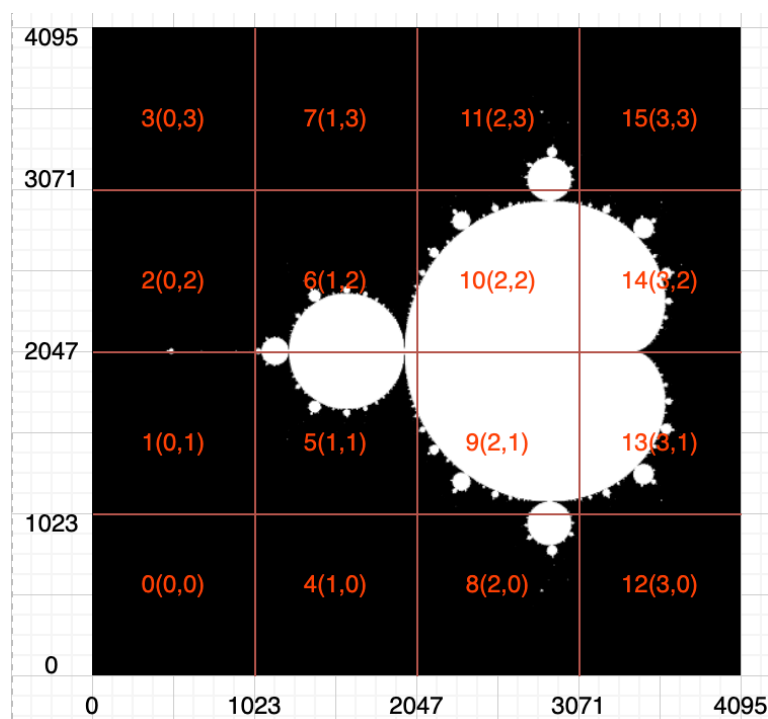


Figure 3: Image from the problem divided between 16 processes. Each process is labelled as rank(x, y). Where rank is the rank of the process in MPI Communicator and (x,y) are the coordinates of the process in the Cartesian Topology. Example 5(1,1); rank=5, x=1 and y=1

This problem requires partitioning an image among multiple processes in a Cartesian Topology. In the [Figure](#) I have partitioned the image into 16 parts for 16 processes.

A partition is described by a **Partition** object, contains a how many processes are there along each dimension of Cartesian Topology. It also contains the co-ordinates of the a given process in that topology.

The work of each processor is described a **Domain** object that is linked to a **Partition** and contains the number of pixels to operate on in each dimension and the start and end index in each dimension. **Domain** uses the co-ordinates of the process from **Partition** object to decide the start and end indices.

Each process works on only one **Domain** object to compute a results and stores them in a result array in row major format. For each process we have `c = malloc((d.nx) * (d.ny) * sizeof(int));` where `c` is the result array and `d` is the local domain for the process.

Each process(except for the root process) sends this result `c` top the root process using `MPI_Send` or others methods like it.

```
1 MPI_Isend(&c[0], d.nx*d.ny, MPI_INT, 0, ping, p.comm, &request);
```

The master process also works on its local domain just like other processes; but it does not send out any information. Also for master process result array `c` is defined as -

```
1 int extrax = IMAGE_WIDTH % p.nx;
2 int extray = IMAGE_HEIGHT % p.ny;
3 c = malloc((d.nx + extrax) * (d.ny + extray) * sizeof(int))
```

where the variables `extrax` and `extray` are the extra space to accommodate for the case where image size along each dimension is not completely divisible by number of processes along that dimension.

As already stated the master process does not send out any information; it is responsible for receiving the results from all other processes and writes them on the image.

It starts with its own `c` and then for each `processi`(non master) -

- Updates its **Partition** object by using the rank of `processi`
- Creates a local domain from this updated **Partition**
- Receives the result from `processi` and stores them in `c`
- Writes the result to image using start and end pixel indices from the **Domain** object

3.1. Results from run_perf.sh

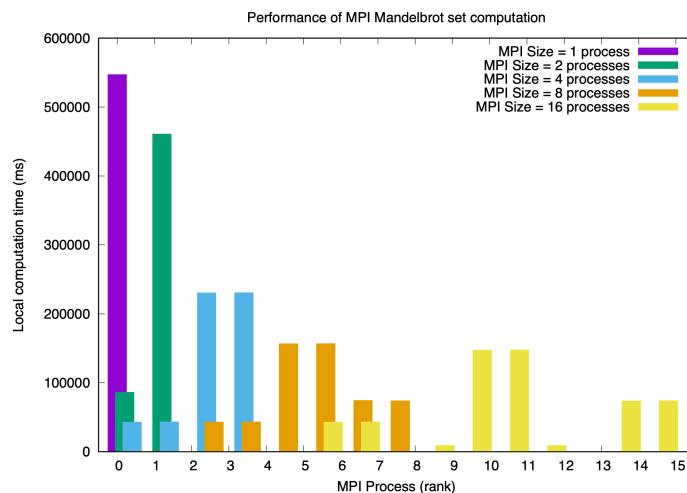


Figure 4: Plots from run_perf.sh

Total Processes	Rank	Local Comp. Time(ms)
16	0	4.377
16	1	109.403
16	2	258.156
16	3	4.297
16	4	8.149
16	5	42940.4
16	6	43109.8
16	7	8.108
16	8	9173.25
16	9	147508
16	10	147640
16	11	9216.06
16	12	10.878
16	13	73820.3
16	14	73882.5
16	15	10.882

Table 1: Table copied from perf.data file generated by run_perf.sh for mandel run with 16 processes

As we can see clearly from the [Table-1](#) and [Figure-3](#) the process whose local domain cover the white pixels do most of the work. This is somehow not reflected in the plot in [Figure-4](#) but can be clearly seen from the table.

4. Option A: Parallel matrix-vector multiplication and the power method [40 Points]

The Power method is an iterative method to find the largest(in absolute value) eigen value of a diagonalizable matrix A , λ and the a vector x , which is the eigen-vector corresponding to λ .

The power iteration algorithm starts with a vector x_0 , which may be an approximation to the dominant eigen-vector or a random vector. The method is described by the recurrence relation

$$x_{k+1} = \frac{Ax_k}{\|x_k\|_2}$$

For this problem matrix multiplication is parallelized by splitting the rows of matrix A among the processes and copying the vector x_k to all the processes such that each process has $\frac{n}{p}$ rows of A ; where n =size of square matrix and p =number of processes Each process computes the $\frac{n}{p}$ elements of x_{k+1} .

The master broadcasts x_k to all the processes using `MPI_Bcast` and collects $\frac{n}{p}$ elements from all process using `MPI_Gather`; this process is repeated many iterations till the power method converges.

4.1. Experiments




Figure 5: Strong Scaling Plot for $n=10000$ and $p=1,2,4,8,16,32$ and 64 . Power method is run for 1000 iterations for each run. (Left)Computation Time vs Processes and (Right)Parallel Efficiency vs Processes




Figure 6: Weak Scaling. Plot of processes vs computation time for $n=\text{int}(\sqrt{p}) * 10000$ and $p=1,2,4,8,16,32$ and 64 . Power method is run for 100 iterations for each run

From [Figure-5](#) it is clear that parallel efficiency decreases sharply after 16 processes as each process is not fully utilized and communication overhead becomes significant. Especially after 16 processes, where computation is faster for 32 processes than 64 processes.

From [Figure-6](#) it can be seen that even though work per process is equal(except at $p=2$ and $p=4$) the trend in computation time is increasing again pointing to the state where inter-process communication overhead becomes significant and affects the computation speed overall.

5. Option B: Parallel PageRank Algorithm and the Power method [40 Points]

6. Task: Quality of the Report [15 Points]

Each project will have 100 points (out of which 15 points will be given to the general quality of the written report).

Additional notes and submission details

Submit the source code files (together with your used **Makefile**) in an archive file (tar, zip, etc.), and summarize your results and the observations for all exercises by writing an extended Latex report. Use the Latex template from the webpage and upload the Latex summary as a PDF to [iCorsi](#).

- Your submission should be a gzipped tar archive, formatted like `project__number__lastname__firstname.zip` or `project__number__lastname__firstname.tgz`. It should contain:
 - all the source codes of your MPI solutions;
 - your write-up with your name `project__number__lastname__firstname.pdf`,
- Submit your .tgz through [Icorsi](#).